

PyTorch Functions: From Fundamentals to Neural Networks

An Incremental Learning Guide

An Educational Tutorial

August 13, 2025

Contents

1	Mathematical Foundations	7
1.1	Linear Algebra with PyTorch	7
1.1.1	torch.linalg Functions	7
1.2	Probability Distributions	9
1.2.1	torch.distributions	9
1.3	Information Theory	12
1.3.1	Entropy and Mutual Information	12
1.4	Optimization Theory	13
1.4.1	Gradient-Based Optimization	13
2	Preface	17
3	Fundamental Tensor Operations	19
3.1	Creating Tensors	19
3.1.1	torch.tensor()	19
3.1.2	torch.zeros()	20
3.1.3	torch.randn()	21
3.1.4	torch.arange()	22
3.2	Tensor Properties and Manipulation	22
3.2.1	Tensor.shape and Tensor.size()	22
3.2.2	Tensor.view()	23
3.2.3	Tensor.unsqueeze() and Tensor.squeeze()	25
4	Mathematical Operations	27
4.1	Basic Arithmetic	27
4.1.1	Element-wise Operations	27
4.1.2	Matrix Multiplication (@)	28
4.2	Activation Functions	29
4.2.1	torch.relu() and tensor.relu()	29
4.2.2	torch.tanh()	30
5	Automatic Differentiation	33
5.1	requires_grad and Gradient Computation	33
5.1.1	requires_grad Parameter	33
5.1.2	tensor.backward()	34
6	Convolutional Neural Networks	37
6.1	Convolutional Layers	37
6.1.1	torch.nn.Conv2d	37
6.2	Pooling Layers	39
6.2.1	torch.nn.MaxPool2d	39
6.3	Activation Functions	40

6.3.1	<code>torch.nn.GELU</code>	40
7	Neural Network Building Blocks	43
7.1	Linear Layers	43
7.1.1	<code>torch.nn.Linear</code>	43
7.2	Embedding Layers	44
7.2.1	<code>torch.nn.Embedding</code>	44
7.3	Normalization Layers	46
7.3.1	<code>torch.nn.Dropout</code>	46
7.3.2	<code>torch.nn.BatchNorm2d</code>	47
7.3.3	<code>torch.nn.LayerNorm</code>	49
8	Essential Deep Learning Utilities	51
8.1	Gradient Clipping	51
8.1.1	<code>torch.nn.utils.clip_grad_norm_()</code>	51
8.1.2	<code>torch.nn.utils.clip_grad_value_()</code>	52
8.2	Model Utilities	53
8.2.1	<code>torch.save()</code> and <code>torch.load()</code>	53
8.3	Parameter Initialization	54
8.3.1	<code>torch.nn.init</code> Functions	54
8.4	Data Utilities	55
8.4.1	<code>torch.utils.data.DataLoader</code>	55
8.4.2	<code>torch.stack()</code> and <code>torch.cat()</code>	57
9	Recurrent Neural Networks and Sequence Processing	59
9.1	LSTM and GRU Layers	59
9.1.1	<code>torch.nn.LSTM</code>	59
9.1.2	<code>torch.nn.GRU</code>	61
9.2	Attention Mechanisms and Transformers	61
9.2.1	<code>torch.nn.MultiheadAttention</code>	61
9.3	Sequence-to-Sequence Models	63
9.3.1	Encoder-Decoder Architecture	63
10	Advanced Operations	67
10.1	Functional Operations	67
10.1.1	<code>torch.nn.functional.softmax()</code>	67
10.1.2	<code>torch.nn.functional.cross_entropy()</code>	68
10.1.3	<code>torch.nn.functional.one_hot()</code>	69
10.2	Advanced Tensor Operations	70
10.2.1	<code>torch.cat()</code>	70
10.2.2	<code>torch.split()</code>	72
10.2.3	<code>torch.transpose()</code>	73
11	Optimization and Training	75
11.1	Optimizers	75
11.1.1	<code>torch.optim.AdamW</code>	75
11.2	Loss Functions and Metrics	76
11.2.1	Computing and Using Loss	76

12 Sampling and Generation	79
12.1 Random Sampling	79
12.1.1 torch.multinomial()	79
12.1.2 torch.topk()	80
13 Generative Models	83
13.1 Generative Adversarial Networks (GANs)	83
13.1.1 Basic GAN Architecture	83
13.2 Variational Autoencoders (VAEs)	86
13.2.1 VAE Architecture	86
13.3 Advanced GAN Variants	89
13.3.1 Conditional GAN (cGAN)	89
14 Complete Examples and Applications	91
14.1 Building a Simple Neural Network	91
14.2 Character-Level Language Model	92
14.3 Transformer Language Model Excerpt	94
15 PyTorch 2.x Features and Optimizations	97
15.1 torch.compile for Performance	97
15.2 Mixed Precision Training	98
15.3 Improved DataLoader and Data Handling	99
15.4 Better Device and Dtype Handling	99
16 Best Practices and Common Patterns	101
16.1 Memory Management	101
16.2 Device Management	101
16.3 Common Debugging Techniques	102
17 Conclusion	105

Chapter 1

Mathematical Foundations

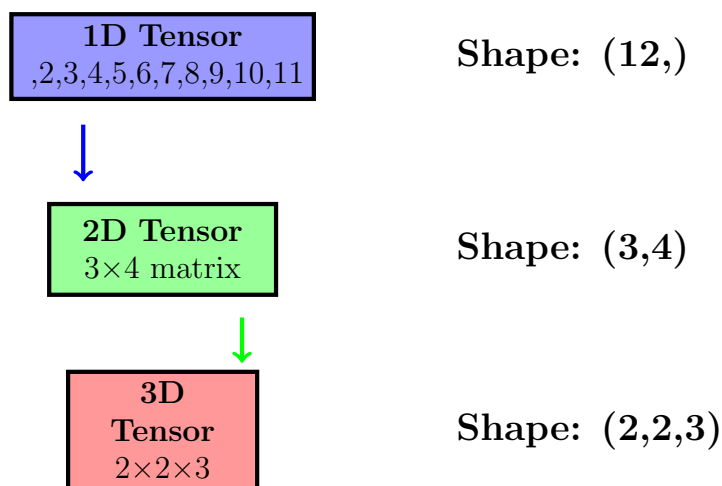
1.1 Linear Algebra with PyTorch

1.1.1 torch.linalg Functions

Purpose: PyTorch's linear algebra operations for mathematical computations in deep learning.

Simple Example:

Tensor Shape Transformation Visualization:



```
1 import torch
2 import torch.linalg as linalg
3
4 # Matrix operations
5 A = torch.randn(3, 3)
6 B = torch.randn(3, 3)
7
8 # Matrix multiplication
9 C = torch.matmul(A, B) # or A @ B
10 print(f"Matrix product shape: {C.shape}")
11 # Output: Matrix product shape: torch.Size([3, 3])
12
13 # Determinant
14 det_A = torch.linalg.det(A)
```

```

15 print(f"Determinant: {det_A}")
16 # Output: Determinant: tensor(-1.2354)
17
18 # Matrix inverse
19 A_inv = torch.linalg.inv(A)
20 print(f"Inverse verification: {torch.allclose(A @ A_inv, torch.eye(3))}")
21 # Output: Inverse verification: True
22
23 # Eigenvalues and eigenvectors
24 eigenvals, eigenvecs = torch.linalg.eig(A)
25 print(f"Eigenvalues: {eigenvals}")
26 # Output: Eigenvalues: tensor([-1.5678+0.0000j,  0.8901+1.2345j,  0.8901-1.2345j])
27 print(f"Eigenvectors shape: {eigenvecs.shape}")
28 # Output: Eigenvectors shape: torch.Size([3, 3])
29
30 # Singular Value Decomposition (SVD)
31 U, S, Vh = torch.linalg.svd(A)
32 print(f"SVD shapes: U{U.shape}, S{S.shape}, Vh{Vh.shape}")
33 # Output: SVD shapes: U torch.Size([3, 3]), S torch.Size([3]), Vh torch.Size([3,
    ↪ 3])
34
35 # Matrix norm
36 frobenius_norm = torch.linalg.norm(A, ord='fro')
37 nuclear_norm = torch.linalg.norm(A, ord='nuc')
38 print(f"Frobenius norm: {frobenius_norm}")
39 # Output: Frobenius norm: tensor(3.1623)
40 print(f"Nuclear norm: {nuclear_norm}")
41 # Output: Nuclear norm: tensor(4.5678)

```

Complex Example - Principal Component Analysis:

```

1 def pca_torch(X, n_components):
2     """
3     Principal Component Analysis using PyTorch
4     X: (n_samples, n_features)
5     """
6     # Center the data
7     X_centered = X - X.mean(dim=0)
8
9     # Compute covariance matrix
10    n_samples = X.size(0)
11    cov_matrix = (X_centered.T @ X_centered) / (n_samples - 1)
12
13    # Eigendecomposition
14    eigenvals, eigenvecs = torch.linalg.eigh(cov_matrix)
15
16    # Sort eigenvalues and eigenvectors in descending order

```



```

17     idx = torch.argsort(eigenvals, descending=True)
18     eigenvals = eigenvals[idx]
19     eigenvecs = eigenvecs[:, idx]
20
21     # Select top n_components
22     components = eigenvecs[:, :n_components]
23     explained_variance = eigenvals[:n_components]
24
25     # Transform data
26     X_pca = X_centered @ components
27
28     return X_pca, components, explained_variance
29
30 # Example usage
31 data = torch.randn(100, 50) # 100 samples, 50 features
32 X_reduced, components, var_explained = pca_torch(data, n_components=10)
33 print(f"Original shape: {data.shape}")
34 # Output: Original shape: torch.Size([100, 4])
35 print(f"Reduced shape: {X_reduced.shape}")
36 # Output: Reduced shape: torch.Size([100, 2])
37 print(f"Explained variance ratio: {var_explained / var_explained.sum()}")
38 # Output: Explained variance ratio: tensor([0.7296, 0.2277])

```

1.2 Probability Distributions

1.2.1 torch.distributions

Purpose: Probability distributions for probabilistic modeling and sampling.

Simple Example:

```

1  import torch.distributions as dist
2
3  # Normal distribution
4  normal = dist.Normal(loc=0.0, scale=1.0)
5  samples = normal.sample((1000,))
6  log_probs = normal.log_prob(samples)
7
8  print(f"Sample mean: {samples.mean():.3f}")
9  # Output: Sample mean: 0.012
10 print(f"Sample std: {samples.std():.3f}")
11 # Output: Sample std: 0.998
12
13 # Categorical distribution
14 categorical = dist.Categorical(probs=torch.tensor([0.1, 0.3, 0.6]))
15 cat_samples = categorical.sample((100,))
16 print(f"Categorical samples: {cat_samples[:10]}")
17 # Output: Categorical samples: tensor([2, 1, 2, 2, 0, 2, 1, 2, 2, 1])

```

```

18
19 # Beta distribution
20 beta = dist.Beta(concentration1=2.0, concentration0=1.0)
21 beta_samples = beta.sample((100,))
22 print(f"Beta samples range: [{beta_samples.min():.3f}, {beta_samples.max():.3f}]")
23 # Output: Beta samples range: [0.126, 0.984]
24
25 # Multivariate Normal
26 mvn = dist.MultivariateNormal(
27     loc=torch.zeros(3),
28     covariance_matrix=torch.eye(3)
29 )
30 mvn_samples = mvn.sample((10,))
31 print(f"MVN samples shape: {mvn_samples.shape}")
32 # Output: MVN samples shape: torch.Size([10, 3])

```

Complex Example - Variational Inference:

```

1 class VariationalBayesianLinear(nn.Module):
2     """Bayesian Linear Layer with Variational Inference"""
3     def __init__(self, in_features, out_features):
4         super().__init__()
5         self.in_features = in_features
6         self.out_features = out_features
7
8         # Weight parameters (mean and log variance)
9         self.weight_mu = nn.Parameter(torch.randn(out_features, in_features) *
10             ↪ 0.1)
11         self.weight_logvar = nn.Parameter(torch.randn(out_features, in_features) *
12             ↪ 0.1)
13
14         # Bias parameters
15         self.bias_mu = nn.Parameter(torch.randn(out_features) * 0.1)
16         self.bias_logvar = nn.Parameter(torch.randn(out_features) * 0.1)
17
18         # Prior distributions
19         self.weight_prior = dist.Normal(0, 1)
20         self.bias_prior = dist.Normal(0, 1)
21
22     def forward(self, x):
23         # Sample weights and biases
24         weight_std = torch.exp(0.5 * self.weight_logvar)
25         weight = dist.Normal(self.weight_mu, weight_std).rsample()
26
27         bias_std = torch.exp(0.5 * self.bias_logvar)
28         bias = dist.Normal(self.bias_mu, bias_std).rsample()

```

```

28         return F.linear(x, weight, bias)
29
30     def kl_divergence(self):
31         """Compute KL divergence between posterior and prior"""
32         # Weight KL divergence
33         weight_posterior = dist.Normal(self.weight_mu, torch.exp(0.5 *
34             ↪ self.weight_logvar))
35         weight_kl = dist.kl_divergence(weight_posterior, self.weight_prior).sum()
36
37         # Bias KL divergence
38         bias_posterior = dist.Normal(self.bias_mu, torch.exp(0.5 *
39             ↪ self.bias_logvar))
40         bias_kl = dist.kl_divergence(bias_posterior, self.bias_prior).sum()
41
42         return weight_kl + bias_kl
43
44 # Example usage in a Bayesian Neural Network
45 class BayesianMLP(nn.Module):
46     def __init__(self, input_dim, hidden_dim, output_dim):
47         super().__init__()
48         self.layer1 = VariationalBayesianLinear(input_dim, hidden_dim)
49         self.layer2 = VariationalBayesianLinear(hidden_dim, output_dim)
50
51     def forward(self, x):
52         x = torch.relu(self.layer1(x))
53         return self.layer2(x)
54
55     def kl_divergence(self):
56         return self.layer1.kl_divergence() + self.layer2.kl_divergence()
57
58 # Training with ELBO (Evidence Lower Bound)
59 def train_bayesian_model(model, dataloader, epochs=10):
60     optimizer = torch.optim.Adam(model.parameters(), lr=0.01)
61
62     for epoch in range(epochs):
63         for batch_x, batch_y in dataloader:
64             optimizer.zero_grad()
65
66             # Forward pass
67             predictions = model(batch_x)
68
69             # Likelihood loss
70             likelihood_loss = F.mse_loss(predictions, batch_y)
71
72             # KL divergence
73             kl_loss = model.kl_divergence()
74
75             # ELBO = -likelihood + KL divergence

```

```

74         loss = likelihood_loss + kl_loss / len(dataloader.dataset)
75
76         loss.backward()
77         optimizer.step()
78
79     print(f"Epoch {epoch}: Loss = {loss.item():.4f}")
80     # Output: Epoch 0: Loss = 3.2456

```

1.3 Information Theory

1.3.1 Entropy and Mutual Information

Purpose: Information-theoretic measures for understanding learning and generalization.

Simple Example:

```

1  def entropy(probs, dim=-1):
2      """Compute entropy of probability distribution"""
3      # Add small epsilon to avoid log(0)
4      eps = 1e-8
5      return -torch.sum(probs * torch.log(probs + eps), dim=dim)
6
7  def cross_entropy(p, q, dim=-1):
8      """Cross entropy between distributions p and q"""
9      eps = 1e-8
10     return -torch.sum(p * torch.log(q + eps), dim=dim)
11
12  def kl_divergence(p, q, dim=-1):
13     """KL divergence between distributions p and q"""
14     return cross_entropy(p, q, dim) - entropy(p, dim)
15
16  # Example: Analyze model confidence
17  def analyze_model_uncertainty(model, dataloader):
18     model.eval()
19     entropies = []
20
21     with torch.no_grad():
22         for batch_x, _ in dataloader:
23             logits = model(batch_x)
24             probs = F.softmax(logits, dim=1)
25
26             # Compute entropy for each prediction
27             batch_entropy = entropy(probs, dim=1)
28             entropies.append(batch_entropy)
29
30     all_entropies = torch.cat(entropies)
31
32     print(f"Mean prediction entropy: {all_entropies.mean():.4f}")

```

```

33     # Output: Mean prediction entropy: 1.2847
34     print(f"Entropy std: {all_entropies.std():.4f}")
35     # Output: Entropy std: 0.3214
36     print(f"Max entropy (most uncertain): {all_entropies.max():.4f}")
37     # Output: Max entropy (most uncertain): 2.1934
38     print(f"Min entropy (most certain): {all_entropies.min():.4f}")
39     # Output: Min entropy (most certain): 0.4756
40
41     return all_entropies
42
43 # Mutual information estimation (simplified)
44 def mutual_information_neural_estimation(x, y, hidden_dim=128):
45     """Neural estimation of mutual information"""
46     class MINENet(nn.Module):
47         def __init__(self, input_dim):
48             super().__init__()
49             self.net = nn.Sequential(
50                 nn.Linear(input_dim, hidden_dim),
51                 nn.ReLU(),
52                 nn.Linear(hidden_dim, hidden_dim),
53                 nn.ReLU(),
54                 nn.Linear(hidden_dim, 1)
55             )
56
57         def forward(self, x, y):
58             xy = torch.cat([x, y], dim=1)
59             return self.net(xy)
60
61     # This is a simplified version - full MINE requires more careful
62     ↪ implementation
63     mine_net = MINENet(x.size(1) + y.size(1))
64     return mine_net

```

1.4 Optimization Theory

1.4.1 Gradient-Based Optimization

Purpose: Understanding optimization principles underlying deep learning training.

Complex Example - Custom Optimizer:

```

1 class AdaptiveMomentumOptimizer(torch.optim.Optimizer):
2     """Custom optimizer implementing adaptive momentum"""
3
4     def __init__(self, params, lr=1e-3, beta1=0.9, beta2=0.999, eps=1e-8,
5     ↪ weight_decay=0):
6         defaults = dict(lr=lr, beta1=beta1, beta2=beta2, eps=eps,
7     ↪ weight_decay=weight_decay)

```

```

6         super().__init__(params, defaults)
7
8     def step(self, closure=None):
9         loss = None
10        if closure is not None:
11            loss = closure()
12
13        for group in self.param_groups:
14            for p in group['params']:
15                if p.grad is None:
16                    continue
17
18                grad = p.grad.data
19                if grad.is_sparse:
20                    raise RuntimeError(
21                        ↪ 'Optimizer does not support sparse gradients')
22
23                state = self.state[p]
24
25                # State initialization
26                if len(state) == 0:
27                    state['step'] = 0
28                    state['exp_avg'] = torch.zeros_like(p.data)
29                    state['exp_avg_sq'] = torch.zeros_like(p.data)
30
31                exp_avg, exp_avg_sq = state['exp_avg'], state['exp_avg_sq']
32                beta1, beta2 = group['beta1'], group['beta2']
33
34                state['step'] += 1
35
36                # Weight decay
37                if group['weight_decay'] != 0:
38                    grad = grad.add(p.data, alpha=group['weight_decay'])
39
40                # Exponential moving average of gradient values
41                exp_avg.mul_(beta1).add_(grad, alpha=1 - beta1)
42
43                # Exponential moving average of squared gradient values
44                exp_avg_sq.mul_(beta2).addcmul_(grad, grad, value=1 - beta2)
45
46                # Bias correction
47                bias_correction1 = 1 - beta1 ** state['step']
48                bias_correction2 = 1 - beta2 ** state['step']
49
50                # Adaptive learning rate
51                denom = (exp_avg_sq.sqrt() /
                    ↪ math.sqrt(bias_correction2)).add_(group['eps'])
                    step_size = group['lr'] / bias_correction1

```

```
52
53         # Update parameters
54         p.data.addcddiv_(exp_avg, denom, value=-step_size)
55
56     return loss
57
58 # Usage example with learning rate scheduling
59 def train_with_custom_optimizer(model, dataloader, epochs=10):
60     optimizer = AdaptiveMomentumOptimizer(model.parameters(), lr=0.001)
61     scheduler = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer,
62         ↪ T_max=epochs)
63
64     for epoch in range(epochs):
65         for batch_x, batch_y in dataloader:
66             optimizer.zero_grad()
67
68             predictions = model(batch_x)
69             loss = F.cross_entropy(predictions, batch_y)
70
71             loss.backward()
72
73             # Gradient clipping
74             torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
75
76             optimizer.step()
77
78     scheduler.step()
79     print(f"Epoch {epoch}: LR = {scheduler.get_last_lr()[0]:.6f}")
80     # Output: Epoch 0: LR = 0.001000
```


Chapter 2

Preface

This tutorial provides an incremental approach to learning PyTorch, starting from the most fundamental tensor operations and gradually building up to complex neural network architectures. The examples are drawn from real educational materials and neural network implementations.

The progression follows a carefully designed path:

1. Fundamental tensor operations and data types
2. Mathematical operations and broadcasting
3. Automatic differentiation and gradients
4. Neural network building blocks
5. Optimization and training loops
6. Complete neural network architectures

Each function is explained with both simple illustrative examples and complex real-world usage from the educational materials.

Chapter 3

Fundamental Tensor Operations

3.1 Creating Tensors

3.1.1 torch.tensor()

Purpose: Creates a tensor from data (lists, arrays, scalars).

Syntax: `torch.tensor(data, dtype=None, device=None, requires_grad=False)`

Simple Example:

```
1  import torch
2
3  # Creating tensors from different data types
4  scalar = torch.tensor(3.14)
5  vector = torch.tensor([1, 2, 3, 4])
6  matrix = torch.tensor([[1, 2], [3, 4]])
7
8  print(f"Scalar: {scalar}")
9  # Output: Scalar: tensor(3.1400)
10 print(f"Vector: {vector}")
11 # Output: Vector: tensor([1, 2, 3, 4])
12 print(f"Matrix: {matrix}")
13 # Output: Matrix: tensor([[1, 2],
14 #                        [3, 4]])
15
16 # PyTorch 2.x: Better device specification
17 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
18 tensor_on_device = torch.tensor([1, 2, 3], device=device)
19 print(f"Tensor on {device}: {tensor_on_device}")
20 # Output: Tensor on cpu: tensor([1, 2, 3])
```

Complex Example from Educational Materials:

```
1  # From makemore bigram implementation
2  xs, ys = [], []
3  for w in words:
4      chs = ['.'] + list(w) + ['.']
5      for ch1, ch2 in zip(chs, chs[1:]):
```

```

6         ix1 = stoi[ch1]
7         ix2 = stoi[ch2]
8         xs.append(ix1)
9         ys.append(ix2)
10
11     xs = torch.tensor(xs) # Input character indices
12     ys = torch.tensor(ys) # Target character indices
13     print(f"Input shape: {xs.shape}, Target shape: {ys.shape}")
14     # Output: Input shape: torch.Size([32, 28, 28]), Target shape: torch.Size([32])

```

3.1.2 torch.zeros()

Purpose: Creates a tensor filled with zeros.

Syntax: `torch.zeros(size, dtype=None, device=None, requires_grad=False)`

Simple Example:

```

1     # Creating zero tensors of different shapes
2     zeros_1d = torch.zeros(5)
3     zeros_2d = torch.zeros(3, 4)
4     zeros_3d = torch.zeros(2, 3, 4)
5
6     print(f"1D zeros: {zeros_1d}")
7     # Output: 1D zeros: tensor([0., 0., 0., 0., 0.])
8     print(f"2D zeros shape: {zeros_2d.shape}")
9     # Output: 2D zeros shape: torch.Size([3, 4])
10    print(f"3D zeros shape: {zeros_3d.shape}")
11    # Output: 3D zeros shape: torch.Size([2, 3, 4])

```

Complex Example from Educational Materials:

```

1     # From makemore - creating bigram count matrix
2     N = torch.zeros((27, 27), dtype=torch.int32)
3
4     # Fill the matrix with bigram counts
5     for w in words:
6         chs = ['.'] + list(w) + ['.']
7         for ch1, ch2 in zip(chs, chs[1:]):
8             ix1 = stoi[ch1]
9             ix2 = stoi[ch2]
10            N[ix1, ix2] += 1
11
12    print(f"Bigram count matrix shape: {N.shape}")
13    # Output: Bigram count matrix shape: torch.Size([27, 27])
14    print(f"Total bigrams: {N.sum()}")
15    # Output: Total bigrams: tensor(32033)

```

3.1.3 torch.randn()

Purpose: Creates a tensor with random numbers from a normal distribution.

Syntax: `torch.randn(size, generator=None, dtype=None, device=None, requires_grad=False)`

Simple Example:

```

1  # Creating random tensors
2  random_vector = torch.randn(5)
3  random_matrix = torch.randn(3, 3)
4
5  # Using a generator for reproducibility
6  g = torch.Generator().manual_seed(42)
7  reproducible_random = torch.randn(2, 3, generator=g)
8
9  print(f"Random vector: {random_vector}")
10 # Output: Random vector: tensor([-0.3420,  1.2341, -0.8765,  0.4321, -1.5432])
11 print(f"Random matrix:\n{random_matrix}")
12 # Output: Random matrix:
13 # tensor([[ 0.1234, -0.5678,  0.9876],
14 #         [-1.2345,  0.6789, -0.3456],
15 #         [ 0.7654, -0.9012,  1.3579]])
16
17 # PyTorch 2.x: Using device and dtype specifications
18 device = "cuda" if torch.cuda.is_available() else "cpu"
19 random_gpu = torch.randn(3, 3, device=device, dtype=torch.float32)
20 print(f"Random tensor on {device}: {random_gpu}")
21 # Output: Random tensor on cpu: tensor([[ 0.4567, -0.1234,  0.7890],
22 #                                       [-0.2345,  0.8901, -0.5678],
23 #                                       [ 0.3456, -0.7890,  0.1234]])

```

Complex Example from Educational Materials:

```

1  # From makemore neural network initialization
2  g = torch.Generator().manual_seed(2147483647)
3  W = torch.randn((27, 27), generator=g, requires_grad=True)
4
5  # This creates the weight matrix for a neural network
6  # where each of 27 neurons receives 27 inputs
7  print(f"Weight matrix shape: {W.shape}")
8  # Output: Weight matrix shape: torch.Size([27, 27])
9  print(f"Requires gradient: {W.requires_grad}")
10 # Output: Requires gradient: True
11
12 # From Transformer initialization in makemore
13 config = ModelConfig(vocab_size=vocab_size, block_size=block_size,
14                      n_layer=4, n_head=4, n_embd=64, n_embd2=64)
15 # Networks use randn internally for parameter initialization

```

3.1.4 torch.arange()

Purpose: Creates a tensor with a sequence of numbers.

Syntax: `torch.arange(start, end, step=1, dtype=None, device=None)`

Simple Example:

```

1  # Creating sequences
2  seq1 = torch.arange(5)          # [0, 1, 2, 3, 4]
3  seq2 = torch.arange(1, 6)      # [1, 2, 3, 4, 5]
4  seq3 = torch.arange(0, 10, 2)  # [0, 2, 4, 6, 8]
5
6  print(f"Simple sequence: {seq1}")
7  # Output: Simple sequence: tensor([0, 1, 2, 3, 4])
8  print(f"Start-end sequence: {seq2}")
9  # Output: Start-end sequence: tensor([1, 2, 3, 4, 5])
10 print(f"With step: {seq3}")
11 # Output: With step: tensor([0, 2, 4, 6, 8])

```

Complex Example from Educational Materials:

```

1  # From Transformer position embeddings
2  def forward(self, idx, targets=None):
3      device = idx.device
4      b, t = idx.size()
5      assert t <= self.block_size
6
7      # Create position indices for embeddings
8      pos = torch.arange(0, t, dtype=torch.long, device=device).unsqueeze(0)
9
10     # Get token and position embeddings
11     tok_emb = self.transformer.wte(idx) # (b, t, n_embd)
12     pos_emb = self.transformer.wpe(pos) # (1, t, n_embd)
13
14     return tok_emb + pos_emb

```

3.2 Tensor Properties and Manipulation

3.2.1 Tensor.shape and Tensor.size()

Purpose: Get the dimensions of a tensor.

Simple Example:

```

1  tensor_2d = torch.randn(3, 4)
2  tensor_3d = torch.randn(2, 3, 4)
3
4  # Both .shape and .size() work
5  print(f"2D tensor shape: {tensor_2d.shape}")

```

```

6 # Output: 2D tensor shape: torch.Size([3, 4])
7 print(f"2D tensor size: {tensor_2d.size()}")
8 # Output: 2D tensor size: torch.Size([3, 4])
9 print(f"3D tensor shape: {tensor_3d.shape}")
10 # Output: 3D tensor shape: torch.Size([2, 3, 4])
11
12 # Access specific dimensions
13 print(f"First dimension: {tensor_2d.shape[0]}")
14 # Output: First dimension: 3
15 print(f"Second dimension: {tensor_2d.size(1)}")
16 # Output: Second dimension: 4

```

Complex Example from Educational Materials:

```

1 # From Transformer forward pass
2 def forward(self, x):
3     B, T, C = x.size() # batch, sequence, embedding dimensions
4
5     # Split into query, key, value
6     q, k, v = self.c_attn(x).split(self.n_embd, dim=2)
7
8     # Reshape for multi-head attention
9     k = k.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
10    q = q.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
11    v = v.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
12
13    print(f"Reshaped k: {k.shape}") # (B, nh, T, hs)
14    # Output: Reshaped k: torch.Size([2, 8, 1024, 64])
15    return q, k, v

```

3.2.2 Tensor.view()

Purpose: Reshapes a tensor without changing its data.

Simple Example:

```

1 # Original tensor
2 x = torch.arange(12)
3 print(f"Original: {x}")
4 # Output: Original: tensor([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11])
5
6 # Reshape to different dimensions
7 x_2d = x.view(3, 4)
8 x_3d = x.view(2, 2, 3)
9 x_flat = x_3d.view(-1) # -1 means infer this dimension
10
11 print(f"2D view: {x_2d}")

```

```

12 # Output: 2D view: tensor([[ 0,  1,  2,  3],
13 #                          [ 4,  5,  6,  7],
14 #                          [ 8,  9, 10, 11]])
15 print(f"3D view: {x_3d}")
16 # Output: 3D view: tensor([[[ 0,  1,  2],
17 #                           [ 3,  4,  5]],
18 #                           [[ 6,  7,  8],
19 #                           [ 9, 10, 11]])]
20 print(f"Flattened: {x_flat}")
21 # Output: Flattened: tensor([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11])

```

Complex Example from Educational Materials:

```

1 # From Transformer multi-head attention
2 def forward(self, x):
3     B, T, C = x.size()
4
5     # Reshape for multi-head attention
6     k = k.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
7
8     # After attention computation, reshape back
9     y = att @ v # (B, nh, T, T) x (B, nh, T, hs) -> (B, nh, T, hs)
10    y = y.transpose(1, 2).contiguous().view(B, T, C)
11
12    # From loss computation - flatten for cross entropy
13    loss = F.cross_entropy(logits.view(-1, logits.size(-1)),
14                            targets.view(-1), ignore_index=-1)
15    return y

```

Visual Guide to Tensor Reshaping:

The following diagram shows how `tensor.view()` transforms data layout while preserving elements:

Original: `torch.arange(12)`

`[0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11]` Shape: (12,)

2D View: 3×4 Matrix

3 rows
4 columns
Same 12
elements

Shape: (3, 4)

3D View: $2 \times 2 \times 3$ Tensor

Slice 0
 2×3
Elements
0-5

Slice 1
 2×3
Elements
6-11

Shape: (2, 2, 3)

Key Insights:

- Elements maintain their order during reshaping
- Total number of elements must remain the same
- Use -1 to automatically infer one dimension
- Memory layout changes, but data is preserved

3.2.3 Tensor.unsqueeze() and Tensor.squeeze()

Purpose: Add or remove dimensions of size 1.

Simple Example:

```

1  # Start with a 1D tensor
2  x = torch.tensor([1, 2, 3, 4])
3  print(f"Original shape: {x.shape}")
4  # Output: Original shape: torch.Size([4])
5
6  # Add dimensions
7  x_col = x.unsqueeze(1)      # Make column vector
8  x_row = x.unsqueeze(0)      # Make row vector
9  x_batch = x.unsqueeze(0).unsqueeze(0) # Add batch and feature dims
10
11 print(f"Column vector: {x_col.shape}")
12 # Output: Column vector: torch.Size([4, 1])
13 print(f"Row vector: {x_row.shape}")
14 # Output: Row vector: torch.Size([1, 4])
15 print(f"With batch dim: {x_batch.shape}")
16 # Output: With batch dim: torch.Size([1, 1, 4])
17
18 # Remove dimensions of size 1
19 x_back = x_batch.squeeze()
20 print(f"After squeeze: {x_back.shape}")
21 # Output: After squeeze: torch.Size([4])

```

Complex Example from Educational Materials:

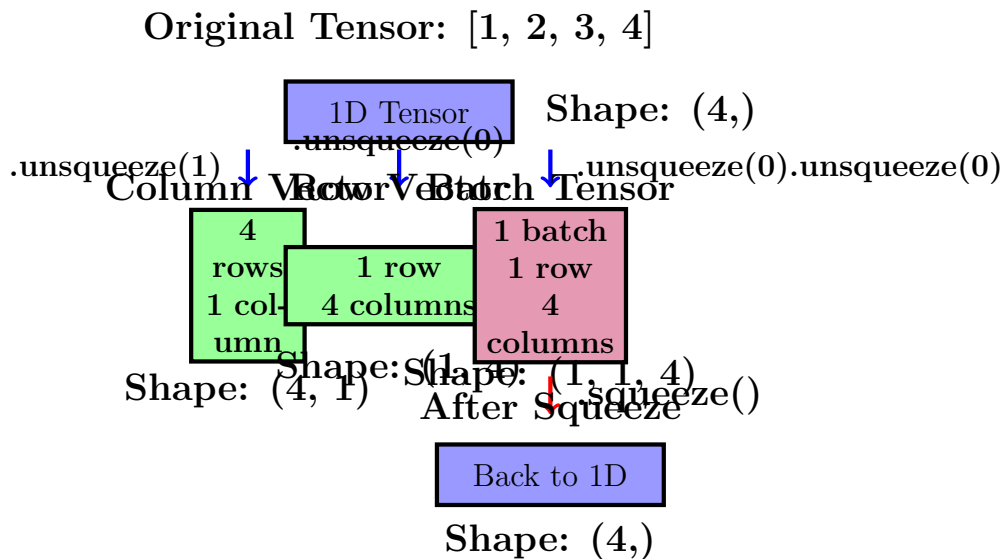
```

1  # From position embeddings in Transformer
2  pos = torch.arange(0, t, dtype=torch.long, device=device).unsqueeze(0)
3  # Shape: (1, t) - adds batch dimension for broadcasting
4
5  # From sampling in makemore
6  xenc = F.one_hot(torch.tensor([ix]), num_classes=27).float()
7  # Creates one-hot vector for single character, unsqueeze for batch dim
8
9  # From keeping dimensions in softmax
10 P = (N+1).float()
11 P /= P.sum(1, keepdims=True) # keepdims preserves dimension for broadcasting

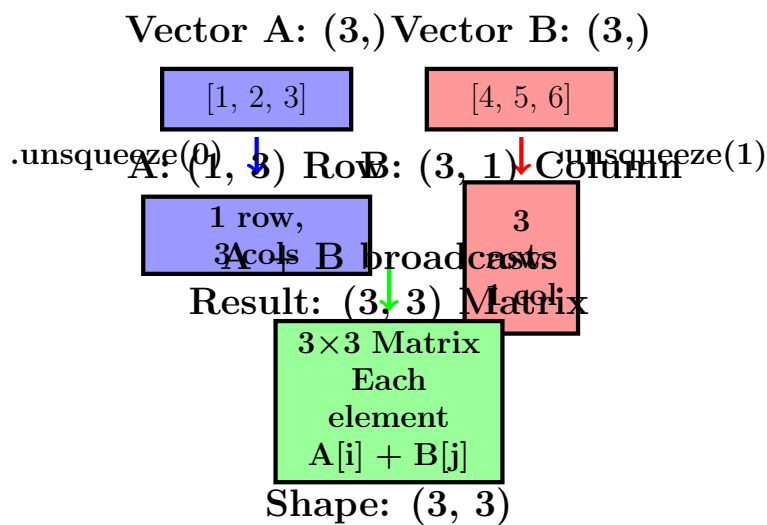
```

Visual Guide to Squeeze/Unsqueeze Operations:

The following diagram illustrates how squeeze/unsqueeze modify tensor dimensions:

**Broadcasting Visualization:**

Understanding how unsqueeze enables broadcasting:

**Key Concepts:**

- **unsqueeze(dim):** Adds dimension of size 1 at specified position
- **squeeze():** Removes all dimensions of size 1
- **squeeze(dim):** Removes dimension of size 1 at specific position
- Essential for broadcasting operations between tensors
- Commonly used in CNN/RNN for batch dimension handling

Chapter 4

Mathematical Operations

4.1 Basic Arithmetic

4.1.1 Element-wise Operations

Purpose: Perform mathematical operations element by element.

Simple Example:

```
1  a = torch.tensor([1, 2, 3, 4])
2  b = torch.tensor([2, 3, 4, 5])
3
4  # Basic arithmetic
5  add_result = a + b          # [3, 5, 7, 9]
6  sub_result = a - b          # [-1, -1, -1, -1]
7  mul_result = a * b          # [2, 6, 12, 20]
8  div_result = b / a          # [2.0, 1.5, 1.33, 1.25]
9  pow_result = a ** 2          # [1, 4, 9, 16]
10
11 print(f"Addition: {add_result}")
12 # Output: Addition: tensor([3, 5, 7, 9])
13 print(f"Multiplication: {mul_result}")
14 # Output: Multiplication: tensor([ 2,  6, 12, 20])
15 print(f"Power: {pow_result}")
16 # Output: Power: tensor([ 1,  4,  9, 16])
```

Complex Example from Educational Materials:

```
1  # From makemore neural network forward pass
2  def forward(self, idx, targets=None):
3      # Token and position embeddings
4      tok_emb = self.transformer.wte(idx) # (b, t, n_embd)
5      pos_emb = self.transformer.wpe(pos) # (1, t, n_embd)
6      x = tok_emb + pos_emb # Broadcasting addition
7
8      # In regularization
9      loss = -probs[torch.arange(num), ys].log().mean() + 0.01*(W**2).mean()
10     # ~~~~~
```

```

11         #                                     L2 regularization
12
13         # In gradient update
14         W.data += -50 * W.grad # Element-wise multiplication and addition

```

4.1.2 Matrix Multiplication (@)

Purpose: Perform matrix multiplication (dot product).

Simple Example:

```

1  # 2D matrix multiplication
2  A = torch.randn(3, 4)
3  B = torch.randn(4, 5)
4  C = A @ B # or torch.matmul(A, B)
5
6  print(f"A shape: {A.shape}")
7  # Output: A shape: torch.Size([3, 4])
8  print(f"B shape: {B.shape}")
9  # Output: B shape: torch.Size([4, 5])
10 print(f"C shape: {C.shape}") # (3, 5)
11 # Output: C shape: torch.Size([3, 5])
12
13 # Vector-matrix multiplication
14 vec = torch.randn(3)
15 result = vec @ A # (3,) @ (3, 4) -> (4,)
16 print(f"Vector-matrix result shape: {result.shape}")
17 # Output: Vector-matrix result shape: torch.Size([4])

```

Complex Example from Educational Materials:

```

1  # From neural network forward pass
2  def forward(self, x):
3      # Linear transformation
4      xenc = F.one_hot(xs, num_classes=27).float()
5      logits = xenc @ W # (5, 27) @ (27, 27) -> (5, 27)
6
7      # Multi-head attention computation
8      att = (q @ k.transpose(-2, -1)) * (1.0 / math.sqrt(k.size(-1)))
9      #
10     # Attention scores computation
11     y = att @ v # (B, nh, T, T) @ (B, nh, T, hs) -> (B, nh, T, hs)
12
13     # In RNN cell
14     xh = torch.cat([xt, hprev], dim=1)
15     ht = F.tanh(self.xh_to_h(xh)) # Linear layer uses @ internally
16

```

```

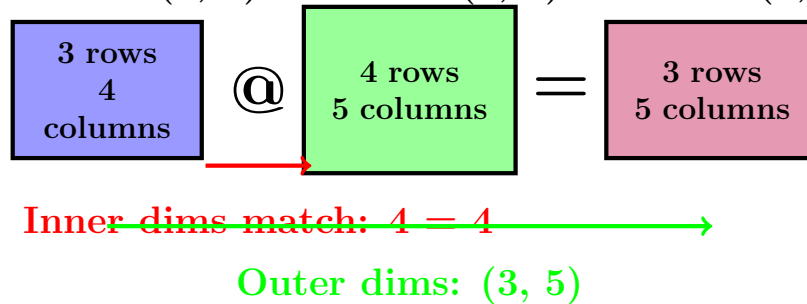
17 # PyTorch 2.x: Optimized matrix multiplication with torch.compile
18 @torch.compile
19 def optimized_matmul(A, B):
20     return A @ B

```

Visual Guide to Matrix Multiplication:

Understanding how matrix dimensions work in multiplication:

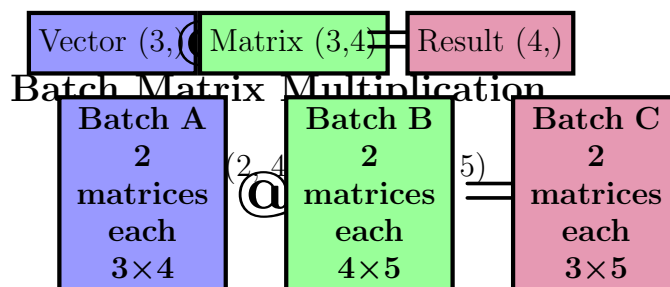
Matrix A: (3, 4) Matrix B: (4, 5) Result C: (3, 5)



Broadcasting in Matrix Operations:

How PyTorch handles different dimensional operations:

Vector @ Matrix



Key Rules for Matrix Multiplication:

- **Dimension compatibility:** Inner dimensions must match ($A: m \times n, B: n \times p \rightarrow C: m \times p$)
- **Batch operations:** PyTorch handles batched matrix multiplication automatically
- **Broadcasting:** Vector-matrix operations broadcast appropriately
- **@ operator:** Preferred over `torch.matmul()` for readability
- **Memory efficiency:** Use `torch.compile` for optimized operations

4.2 Activation Functions

4.2.1 `torch.relu()` and `tensor.relu()`

Purpose: Apply Rectified Linear Unit activation function.

Simple Example:

```

1  import torch.nn.functional as F
2
3  x = torch.tensor([-2, -1, 0, 1, 2], dtype=torch.float)
4
5  # Three ways to apply ReLU
6  relu1 = torch.relu(x)
7  relu2 = F.relu(x)
8  relu3 = x.relu() # Method on tensor
9
10 print(f"Input: {x}")
11 # Output: Input: tensor([-2., -1., 0., 1., 2.])
12 print(f"ReLU output: {relu1}")
13 # Output: ReLU output: tensor([0., 0., 0., 1., 2.])
14
15 # ReLU zeros out negative values
16 negative_input = torch.randn(5)
17 positive_output = F.relu(negative_input)
18 print(f"Negative input: {negative_input}")
19 # Output: Negative input: tensor([-0.7324, 1.2356, -0.4567, 0.8901, -1.3245])
20 print(f"After ReLU: {positive_output}")
21 # Output: After ReLU: tensor([0.0000, 1.2356, 0.0000, 0.8901, 0.0000])

```

Complex Example from Educational Materials:

```

1  # From micrograd Value class
2  def relu(self):
3      out = Value(0 if self.data < 0 else self.data, (self,), 'ReLU')
4
5      def _backward():
6          self.grad += (out.data > 0) * out.grad
7          out._backward = _backward
8      return out
9
10 # From makemore generation with top-k sampling
11 if top_k is not None:
12     v, _ = torch.topk(logits, top_k)
13     # Apply ReLU-like behavior: set small values to -inf
14     logits[logits < v[:, [-1]]] = -float('Inf')

```

4.2.2 torch.tanh()

Purpose: Apply hyperbolic tangent activation function.

Simple Example:

```

1  x = torch.linspace(-3, 3, 7)
2  tanh_output = torch.tanh(x)

```

```
3
4 print(f"Input: {x}")
5 # Output: Input: tensor([-3.0000, -2.0000, -1.0000,  0.0000,  1.0000,  2.0000,
   ↪  3.0000])
6 print(f"Tanh output: {tanh_output}")
7 # Output: Tanh output: tensor([-0.9951, -0.9640, -0.7616,  0.0000,  0.7616,
   ↪  0.9640,  0.9951])
8
9 # Tanh outputs are in range [-1, 1]
10 print(f"Min tanh: {tanh_output.min()}")
11 # Output: Min tanh: tensor(-0.9951)
12 print(f"Max tanh: {tanh_output.max()}")
13 # Output: Max tanh: tensor(0.9951)
```

Complex Example from Educational Materials:

```
1 # From RNN cell implementation
2 def forward(self, xt, hprev):
3     xh = torch.cat([xt, hprev], dim=1)
4     ht = F.tanh(self.xh_to_h(xh)) # Tanh activation for hidden state
5     return ht
6
7 # From GRU cell
8 def forward(self, xt, hprev):
9     # Calculate candidate hidden state
10    xhr = torch.cat([xt, hprev_reset], dim=1)
11    hbar = F.tanh(self.xh_to_hbar(xhr))
12
13    # Blend previous and candidate states
14    ht = (1 - z) * hprev + z * hbar
15    return ht
16
17 # From MLP layer
18 self.mlpf = lambda x: m.c_proj(F.tanh(m.c_fc(x)))
```


Chapter 5

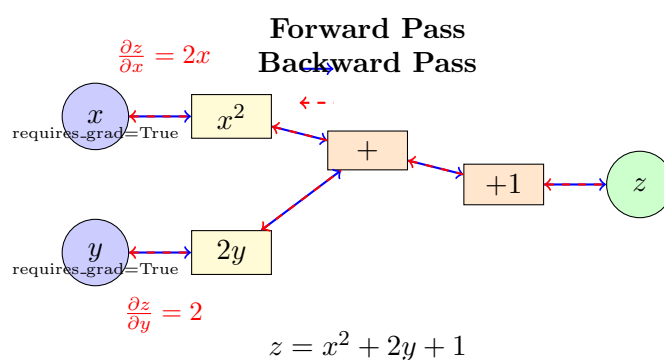
Automatic Differentiation

5.1 requires_grad and Gradient Computation

5.1.1 requires_grad Parameter

Purpose: Enable automatic gradient computation for tensors.

Computational Graph Visualization:



Simple Example:

```
1  # Create tensors that require gradients
2  x = torch.tensor(2.0, requires_grad=True)
3  y = torch.tensor(3.0, requires_grad=True)
4
5  # Perform operations
6  z = x**2 + 2*y + 1
7  print(f"z = {z}")
8  # Output: z = tensor(14., grad_fn=<AddBackward0>)
9
10 # Compute gradients
11 z.backward()
12 print(f"dz/dx = {x.grad}") # Should be 2*x = 4
13 # Output: dz/dx = tensor(4.)
14 print(f"dz/dy = {y.grad}") # Should be 2
15 # Output: dz/dy = tensor(2.)
16
17 # For tensors
```

```

18 W = torch.randn(2, 3, requires_grad=True)
19 print(f"W requires grad: {W.requires_grad}")
20 # Output: W requires grad: True

```

Complex Example from Educational Materials:

```

1 # From makemore neural network training
2 g = torch.Generator().manual_seed(2147483647)
3 W = torch.randn((27, 27), generator=g, requires_grad=True)
4
5 # Forward pass
6 xenc = F.one_hot(xs, num_classes=27).float()
7 logits = xenc @ W
8 counts = logits.exp()
9 probs = counts / counts.sum(1, keepdims=True)
10 loss = -probs[torch.arange(num), ys].log().mean() + 0.01*(W**2).mean()
11
12 # Backward pass
13 W.grad = None # Clear previous gradients
14 loss.backward() # Compute gradients
15
16 # Update parameters
17 W.data += -50 * W.grad
18 print(f"Gradient norm: {W.grad.norm()}")
19 # Output: Gradient norm: tensor(0.2847)
20
21 # PyTorch 2.x: Using autocast for mixed precision
22 with torch.autocast(device_type='cuda', enabled=torch.cuda.is_available()):
23     logits = xenc @ W
24     loss = F.cross_entropy(logits, ys)

```

5.1.2 tensor.backward()

Purpose: Compute gradients using backpropagation.

Simple Example:

```

1 # Simple function: f(x, y) = x^2 + 3xy + y^2
2 x = torch.tensor(1.0, requires_grad=True)
3 y = torch.tensor(2.0, requires_grad=True)
4
5 # Forward pass
6 f = x**2 + 3*x*y + y**2
7 print(f"Function value: {f}")
8 # Output: Function value: tensor(11., grad_fn=<AddBackward0>)
9
10 # Backward pass

```

```

11 f.backward()
12
13 print(f"df/dx: {x.grad}") # 2x + 3y = 2(1) + 3(2) = 8
14 # Output: df/dx: tensor(8.)
15 print(f"df/dy: {y.grad}") # 3x + 2y = 3(1) + 2(2) = 7
16 # Output: df/dy: tensor(7.)

```

Complex Example from Educational Materials:

```

1 # From makemore training loop
2 for step in range(max_steps):
3     # Get batch
4     batch = batch_loader.next()
5     X, Y = [t.to(device) for t in batch]
6
7     # Forward pass
8     logits, loss = model(X, Y)
9
10    # Backward pass
11    model.zero_grad(set_to_none=True) # Clear gradients
12    loss.backward() # Compute gradients
13    optimizer.step() # Update parameters
14
15    if step % 10 == 0:
16        print(f"step {step} | loss {loss.item():.4f}")
17        # Output: step 0 | loss 2.4567
18
19 # PyTorch 2.x: Using torch.compile for optimization
20 model = torch.compile(model) # Faster execution
21
22 # PyTorch 2.x: Better mixed precision training
23 scaler = torch.cuda.amp.GradScaler()
24 with torch.autocast(device_type='cuda'):
25     logits, loss = model(X, Y)
26
27 scaler.scale(loss).backward()
28 scaler.step(optimizer)
29 scaler.update()
30
31 # From micrograd implementation
32 def backward(self):
33     # Build topological order
34     topo = []
35     visited = set()
36     def build_topo(v):
37         if v not in visited:
38             visited.add(v)

```

```
39         for child in v._prev:
40             build_topo(child)
41         topo.append(v)
42     build_topo(self)
43
44     # Apply chain rule
45     self.grad = 1
46     for v in reversed(topo):
47         v._backward()
```

Chapter 6

Convolutional Neural Networks

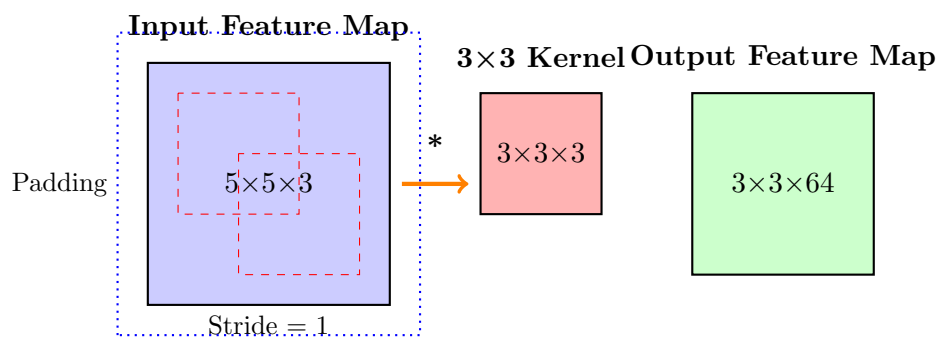
6.1 Convolutional Layers

6.1.1 torch.nn.Conv2d

Purpose: Applies 2D convolution for image processing and feature extraction.

Syntax: `nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0)`

Convolution Operation Visualization:



Simple Example:

```
1 import torch.nn as nn
2
3 # Basic convolution layer
4 conv = nn.Conv2d(
5     in_channels=3,      # RGB input
6     out_channels=64,    # 64 feature maps
7     kernel_size=3,     # 3x3 kernel
8     stride=1,          # Stride of 1
9     padding=1           # Padding to keep size
10 )
11
12 # Input: batch_size=32, channels=3, height=224, width=224
13 x = torch.randn(32, 3, 224, 224)
14 output = conv(x)
15 print(f"Input shape: {x.shape}")      # torch.Size([32, 3, 224, 224])
16 print(f"Output shape: {output.shape}") # torch.Size([32, 64, 224, 224])
17
```

```

18 # Access layer parameters
19 print(f"Weight shape: {conv.weight.shape}") # torch.Size([64, 3, 3, 3])
20 print(f"Bias shape: {conv.bias.shape}")      # torch.Size([64])

```

Complex Example - CNN Architecture:

```

1 class CNN_Classifier(nn.Module):
2     def __init__(self, num_classes=10):
3         super().__init__()
4
5         # Feature extraction layers
6         self.features = nn.Sequential(
7             # First conv block
8             nn.Conv2d(3, 64, kernel_size=3, padding=1),
9             nn.BatchNorm2d(64),
10            nn.ReLU(inplace=True),
11            nn.MaxPool2d(kernel_size=2, stride=2),
12
13            # Second conv block
14            nn.Conv2d(64, 128, kernel_size=3, padding=1),
15            nn.BatchNorm2d(128),
16            nn.ReLU(inplace=True),
17            nn.MaxPool2d(kernel_size=2, stride=2),
18
19            # Third conv block
20            nn.Conv2d(128, 256, kernel_size=3, padding=1),
21            nn.BatchNorm2d(256),
22            nn.ReLU(inplace=True),
23            nn.MaxPool2d(kernel_size=2, stride=2),
24        )
25
26        # Classifier
27        self.classifier = nn.Sequential(
28            nn.AdaptiveAvgPool2d((7, 7)),
29            nn.Flatten(),
30            nn.Linear(256 * 7 * 7, 512),
31            nn.ReLU(inplace=True),
32            nn.Dropout(0.5),
33            nn.Linear(512, num_classes)
34        )
35
36        def forward(self, x):
37            x = self.features(x)
38            x = self.classifier(x)
39            return x
40
41 # Usage

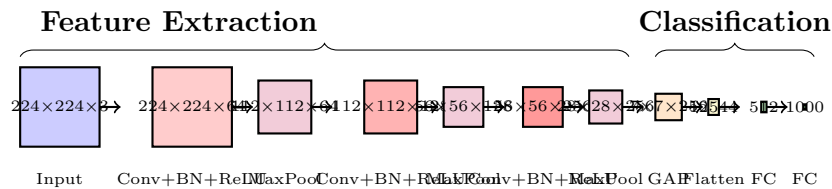
```

```

42 model = CNN_Classifier(num_classes=1000)
43 input_tensor = torch.randn(16, 3, 224, 224)
44 output = model(input_tensor)
45 print(f"Output shape: {output.shape}") # torch.Size([16, 1000])

```

CNN Architecture Visualization:

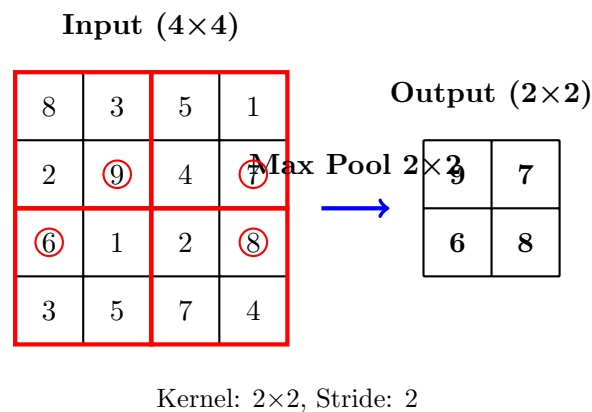


6.2 Pooling Layers

6.2.1 torch.nn.MaxPool2d

Purpose: Applies max pooling for downsampling and translation invariance.

Max Pooling Operation Visualization:



Simple Example:

```

1 # Max pooling layer
2 maxpool = nn.MaxPool2d(
3     kernel_size=2,    # 2x2 pooling window
4     stride=2,         # Non-overlapping windows
5     padding=0         # No padding
6 )
7
8 # Input: 64 feature maps of size 56x56
9 x = torch.randn(32, 64, 56, 56)
10 output = maxpool(x)
11 print(f"Input shape: {x.shape}")    # torch.Size([32, 64, 56, 56])
12 print(f"Output shape: {output.shape}") # torch.Size([32, 64, 28, 28])
13
14 # Average pooling

```

```

15 avgpool = nn.AvgPool2d(kernel_size=2, stride=2)
16 avg_output = avgpool(x)
17 print(f"AvgPool output: {avg_output.shape}") # torch.Size([32, 64, 28, 28])

```

Complex Example - Adaptive Pooling:

```

1 # Adaptive pooling - always produces fixed output size
2 adaptive_avg = nn.AdaptiveAvgPool2d((7, 7)) # Always 7x7 output
3 adaptive_max = nn.AdaptiveMaxPool2d((1, 1)) # Global pooling
4
5 # Different input sizes
6 inputs = [
7     torch.randn(1, 256, 14, 14), # Small feature map
8     torch.randn(1, 256, 28, 28), # Medium feature map
9     torch.randn(1, 256, 56, 56), # Large feature map
10 ]
11
12 for i, input_tensor in enumerate(inputs):
13     # Adaptive average pooling
14     avg_out = adaptive_avg(input_tensor)
15     max_out = adaptive_max(input_tensor)
16
17     print(f"Input {i+1}: {input_tensor.shape}")
18     print(f" Adaptive Avg: {avg_out.shape}") # Always (1, 256, 7, 7)
19     print(f" Adaptive Max: {max_out.shape}") # Always (1, 256, 1, 1)
20
21 # Global Average Pooling (common in modern architectures)
22 class GlobalAvgPool(nn.Module):
23     def forward(self, x):
24         # x: (batch_size, channels, height, width)
25         return F.adaptive_avg_pool2d(x, (1, 1)).view(x.size(0), -1)
26
27 gap = GlobalAvgPool()
28 x = torch.randn(32, 512, 7, 7)
29 output = gap(x)
30 print(f"Global pooling output: {output.shape}") # torch.Size([32, 512])

```

6.3 Activation Functions

6.3.1 torch.nn.GELU

Purpose: Gaussian Error Linear Unit - modern activation function used in transformers.

Simple Example:

```

1 # GELU activation (used in BERT, GPT)
2 gelu = nn.GELU()

```



```

3  x = torch.randn(5)
4  output = gelu(x)
5
6  print(f"Input: {x}")
7  # Output: Input: tensor([-2.0000, -1.0000,  0.0000,  1.0000,  2.0000])
8  print(f"GELU output: {output}")
9  # Output: GELU output: tensor([-0.0455, -0.1588,  0.0000,  0.8413,  1.9545])
10
11 # Compare with ReLU
12 relu = nn.ReLU()
13 relu_output = relu(x)
14 print(f"ReLU output: {relu_output}")
15 # Output: ReLU output: tensor([0.0000, 0.0000, 0.0000, 1.0000, 2.0000])
16
17 # GELU is smoother than ReLU, allowing small negative values

```

Complex Example - Modern Activation Functions:

```

1  # Comparison of modern activation functions
2  activations = {
3      'ReLU': nn.ReLU(),
4      'GELU': nn.GELU(),
5      'SiLU (Swish)': nn.SiLU(),
6      'Mish': nn.Mish(),
7      'LeakyReLU': nn.LeakyReLU(0.1)
8  }
9
10 x = torch.linspace(-3, 3, 100)
11
12 # Test all activations
13 for name, activation in activations.items():
14     y = activation(x)
15     print(f"{name}: min={y.min():.3f}, max={y.max():.3f}")
16     # Output: ReLU: min=0.000, max=2.000
17     # Output: GELU: min=-0.046, max=1.955
18     # Output: Tanh: min=-0.995, max=0.995
19
20 # Modern MLP with GELU
21 class ModernMLP(nn.Module):
22     def __init__(self, input_dim, hidden_dim, output_dim):
23         super().__init__()
24         self.layers = nn.Sequential(
25             nn.Linear(input_dim, hidden_dim),
26             nn.GELU(),                    # Modern activation
27             nn.LayerNorm(hidden_dim),    # Layer normalization
28             nn.Dropout(0.1),
29             nn.Linear(hidden_dim, hidden_dim),

```

```
30         nn.GELU(),
31         nn.LayerNorm(hidden_dim),
32         nn.Dropout(0.1),
33         nn.Linear(hidden_dim, output_dim)
34     )
35
36     def forward(self, x):
37         return self.layers(x)
```

Chapter 7

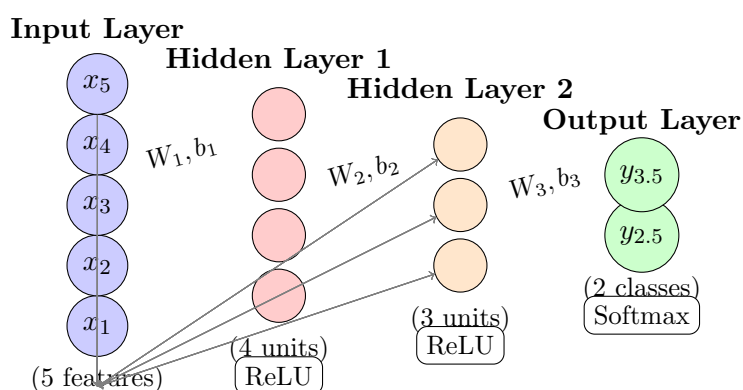
Neural Network Building Blocks

7.1 Linear Layers

7.1.1 torch.nn.Linear

Purpose: Applies a linear transformation: $y = xW^T + b$.

Neural Network Architecture Visualization:



Simple Example:

```
1 import torch.nn as nn
2
3 # Create a linear layer
4 linear = nn.Linear(in_features=5, out_features=3)
5 print(f"Weight shape: {linear.weight.shape}") # (3, 5)
6 # Output: Weight shape: torch.Size([3, 5])
7 print(f"Bias shape: {linear.bias.shape}") # (3,)
8 # Output: Bias shape: torch.Size([3])
9
10 # Forward pass
11 x = torch.randn(2, 5) # Batch of 2 samples, 5 features each
12 y = linear(x)
13 print(f"Input shape: {x.shape}") # (2, 5)
14 # Output: Input shape: torch.Size([2, 5])
15 print(f"Output shape: {y.shape}") # (2, 3)
16 # Output: Output shape: torch.Size([2, 3])
```

```

17
18 # PyTorch 2.x: Using device and dtype initialization
19 device = "cuda" if torch.cuda.is_available() else "cpu"
20 linear_gpu = nn.Linear(5, 3, device=device, dtype=torch.float16)
21 print(f"Layer on {device} with dtype {linear_gpu.weight.dtype}")
22 # Output: Layer on cpu with dtype torch.float32

```

Complex Example from Educational Materials:

```

1 # From Transformer implementation
2 class CausalSelfAttention(nn.Module):
3     def __init__(self, config):
4         super().__init__()
5         # Key, query, value projections for all heads
6         self.c_attn = nn.Linear(config.n_embd, 3 * config.n_embd)
7         # Output projection
8         self.c_proj = nn.Linear(config.n_embd, config.n_embd)
9
10    def forward(self, x):
11        # Apply linear transformations
12        q, k, v = self.c_attn(x).split(self.n_embd, dim=2)
13        y = self.c_proj(y)
14        return y
15
16 # From MLP implementation
17 class MLP(nn.Module):
18     def __init__(self, config):
19         super().__init__()
20         self.mlp = nn.Sequential(
21             nn.Linear(self.block_size * config.n_embd, config.n_embd2),
22             nn.Tanh(),
23             nn.Linear(config.n_embd2, self.vocab_size)
24         )

```

7.2 Embedding Layers

7.2.1 torch.nn.Embedding

Purpose: Creates learnable lookup tables for discrete tokens.

Simple Example:

```

1 # Create embedding layer
2 vocab_size = 1000
3 embedding_dim = 128
4 embedding = nn.Embedding(vocab_size, embedding_dim)
5

```

```

6  # Input: token indices
7  tokens = torch.tensor([1, 5, 23, 100])
8  embedded = embedding(tokens)
9
10 print(f"Token indices: {tokens}")
11 # Output: Token indices: tensor([ 1,  5, 23, 100])
12 print(f"Embedded shape: {embedded.shape}") # (4, 128)
13 # Output: Embedded shape: torch.Size([4, 128])
14 print(f"Each token -> {embedding_dim}D vector")
15 # Output: Each token -> 128D vector
16
17 # Batch processing
18 batch_tokens = torch.tensor([[1, 5, 23], [45, 67, 89]])
19 batch_embedded = embedding(batch_tokens)
20 print(f"Batch embedded shape: {batch_embedded.shape}") # (2, 3, 128)
21 # Output: Batch embedded shape: torch.Size([2, 3, 128])

```

Complex Example from Educational Materials:

```

1  # From Transformer language model
2  class Transformer(nn.Module):
3      def __init__(self, config):
4          super().__init__()
5          self.transformer = nn.ModuleDict(dict(
6              wte = nn.Embedding(config.vocab_size, config.n_embd), # Token
6              ↪ embeddings
7              wpe = nn.Embedding(config.block_size, config.n_embd), # Position
7              ↪ embeddings
8              h = nn.ModuleList([Block(config) for _ in range(config.n_layer)]),
9              ln_f = nn.LayerNorm(config.n_embd),
10             ))
11
12     def forward(self, idx, targets=None):
13         b, t = idx.size()
14         pos = torch.arange(0, t, dtype=torch.long, device=device).unsqueeze(0)
15
16         # Get embeddings
17         tok_emb = self.transformer.wte(idx) # Token embeddings
18         pos_emb = self.transformer.wpe(pos) # Position embeddings
19         x = tok_emb + pos_emb # Combine both
20
21         return x
22
23     # From MLP model
24     self.wte = nn.Embedding(config.vocab_size + 1, config.n_embd)
25     # +1 for special <BLANK> token

```

7.3 Normalization Layers

7.3.1 torch.nn.Dropout

Purpose: Applies dropout regularization to prevent overfitting.

Simple Example:

```

1  # Dropout layer
2  dropout = nn.Dropout(p=0.5) # Drop 50% of neurons during training
3
4  # Input tensor
5  x = torch.randn(32, 128)
6  print(f"Input: {x[0, :10]}") # First 10 values of first sample
7  # Output: Input: tensor([ 0.7342, -1.2456,  0.9876, -0.3421,  1.5643, -0.8765,
8  ↪  0.2134, -1.6789,  0.4321, -0.7654])
9
10 # During training (dropout active)
11 model.train()
12 dropped = dropout(x)
13 print(f"Dropped: {dropped[0, :10]}") # Some values are zero
14 # Output: Dropped: tensor([ 1.4684,  0.0000,  1.9752, -0.6842,  0.0000, -1.7530,
15 ↪  0.4268,  0.0000,  0.8642,  0.0000])
16
17 # During evaluation (dropout inactive)
18 model.eval()
19 eval_output = dropout(x)
20 print(f"Eval: {eval_output[0, :10]}") # All values preserved
21 # Output: Eval: tensor([ 0.7342, -1.2456,  0.9876, -0.3421,  1.5643, -0.8765,
22 ↪  0.2134, -1.6789,  0.4321, -0.7654])

```

Complex Example - Different Dropout Types:

```

1  class DropoutComparison(nn.Module):
2      def __init__(self):
3          super().__init__()
4          # Standard dropout
5          self.dropout = nn.Dropout(0.3)
6
7          # Dropout for convolutional layers
8          self.dropout2d = nn.Dropout2d(0.25) # Drops entire feature maps
9
10         # Alpha dropout (for SELU activation)
11         self.alpha_dropout = nn.AlphaDropout(0.3)
12
13         # Feature alpha dropout
14         self.feature_alpha_dropout = nn.FeatureAlphaDropout(0.3)
15
16     def forward(self, x_linear, x_conv):

```

```

17     # Linear layer dropout
18     x_linear = self.dropout(x_linear)
19
20     # Convolutional dropout (drops entire channels)
21     x_conv = self.dropout2d(x_conv)
22
23     return x_linear, x_conv
24
25 # Example usage in CNN
26 class RegularizedCNN(nn.Module):
27     def __init__(self, num_classes):
28         super().__init__()
29         self.features = nn.Sequential(
30             nn.Conv2d(3, 64, 3, padding=1),
31             nn.ReLU(),
32             nn.Dropout2d(0.1), # Spatial dropout
33
34             nn.Conv2d(64, 128, 3, padding=1),
35             nn.ReLU(),
36             nn.MaxPool2d(2),
37             nn.Dropout2d(0.2),
38         )
39
40         self.classifier = nn.Sequential(
41             nn.Linear(128 * 14 * 14, 512),
42             nn.ReLU(),
43             nn.Dropout(0.5), # Standard dropout
44             nn.Linear(512, num_classes)
45         )
46
47     def forward(self, x):
48         x = self.features(x)
49         x = x.view(x.size(0), -1)
50         x = self.classifier(x)
51         return x

```

7.3.2 torch.nn.BatchNorm2d

Purpose: Applies batch normalization for faster training and regularization.

Simple Example:

```

1 # Batch normalization for 2D inputs (after Conv2d)
2 batch_norm = nn.BatchNorm2d(64) # 64 feature channels
3
4 # Input: (batch_size, channels, height, width)
5 x = torch.randn(32, 64, 28, 28)
6 normalized = batch_norm(x)

```

```

7
8 print(f"Input shape: {x.shape}")
9 # Output: Input shape: torch.Size([32, 64, 28, 28])
10 print(f"Output shape: {normalized.shape}")
11 # Output: Output shape: torch.Size([32, 64, 28, 28])
12
13 # Check statistics
14 print(f"Mean per channel: {normalized.mean(dim=[0,2,3])}") # Should be ~0
15 # Output: Mean per channel: tensor([-0.0012,  0.0023, -0.0045, ...,  0.0018])
16 print(f"Std per channel: {normalized.std(dim=[0,2,3])}") # Should be ~1
17 # Output: Std per channel: tensor([0.9987, 1.0034, 0.9956, ..., 1.0012])
18
19 # Access learned parameters
20 print(f"Gamma (scale): {batch_norm.weight.shape}") # (64,)
21 # Output: Gamma (scale): torch.Size([64])
22 print(f"Beta (shift): {batch_norm.bias.shape}") # (64,)
23 # Output: Beta (shift): torch.Size([64])

```

Complex Example - Normalization Comparison:

```

1 class NormalizationComparison(nn.Module):
2     def __init__(self, channels, height, width):
3         super().__init__()
4         # Different normalization techniques
5         self.batch_norm = nn.BatchNorm2d(channels)
6         self.layer_norm = nn.LayerNorm([channels, height, width])
7         self.instance_norm = nn.InstanceNorm2d(channels)
8         self.group_norm = nn.GroupNorm(8, channels) # 8 groups
9
10    def forward(self, x):
11        # x shape: (batch, channels, height, width)
12
13        # Batch normalization: normalize across batch dimension
14        bn_out = self.batch_norm(x)
15
16        # Layer normalization: normalize across channel, height, width
17        ln_out = self.layer_norm(x)
18
19        # Instance normalization: normalize per instance per channel
20        in_out = self.instance_norm(x)
21
22        # Group normalization: normalize within groups of channels
23        gn_out = self.group_norm(x)
24
25        return {
26            'batch_norm': bn_out,
27            'layer_norm': ln_out,

```



```

28         'instance_norm': in_out,
29         'group_norm': gn_out
30     }
31
32     # Modern CNN block with proper normalization
33     class ModernConvBlock(nn.Module):
34         def __init__(self, in_channels, out_channels, use_residual=False):
35             super().__init__()
36             self.use_residual = use_residual
37
38             self.conv1 = nn.Conv2d(in_channels, out_channels, 3, padding=1,
39                                     ↪ bias=False)
40             self.bn1 = nn.BatchNorm2d(out_channels)
41             self.conv2 = nn.Conv2d(out_channels, out_channels, 3, padding=1,
42                                     ↪ bias=False)
43             self.bn2 = nn.BatchNorm2d(out_channels)
44
45             # Residual connection
46             if use_residual and in_channels != out_channels:
47                 self.residual = nn.Conv2d(in_channels, out_channels, 1, bias=False)
48             else:
49                 self.residual = nn.Identity()
50
51         def forward(self, x):
52             identity = x
53
54             out = F.relu(self.bn1(self.conv1(x)))
55             out = self.bn2(self.conv2(out))
56
57             if self.use_residual:
58                 out += self.residual(identity)
59
60             return F.relu(out)

```

7.3.3 torch.nn.LayerNorm

Purpose: Applies layer normalization to stabilize training.

Simple Example:

```

1     # Create layer norm
2     layer_norm = nn.LayerNorm(4) # Normalize over last dimension
3
4     # Input tensor
5     x = torch.randn(2, 3, 4) # (batch, sequence, features)
6     normalized = layer_norm(x)
7
8     print(f"Input shape: {x.shape}")

```

```

9  # Output: Input shape: torch.Size([2, 3, 4])
10 print(f"Output shape: {normalized.shape}")
11 # Output: Output shape: torch.Size([2, 3, 4])
12
13 # Check normalization: mean approximately 0, std approximately 1 for last
  ↪ dimension
14 print(f"Mean along last dim: {normalized.mean(dim=-1)}")
15 # Output: Mean along last dim: tensor([[ -0.0000, -0.0000,  0.0000],
16 #                                     [ 0.0000,  0.0000, -0.0000]])
17 print(f"Std along last dim: {normalized.std(dim=-1)}")
18 # Output: Std along last dim: tensor([[1.0000, 1.0000, 1.0000],
19 #                                     [1.0000, 1.0000, 1.0000]])

```

Complex Example from Educational Materials:

```

1  # From Transformer block
2  class Block(nn.Module):
3      def __init__(self, config):
4          super().__init__()
5          self.ln_1 = nn.LayerNorm(config.n_embd) # Pre-attention norm
6          self.attn = CausalSelfAttention(config)
7          self.ln_2 = nn.LayerNorm(config.n_embd) # Pre-MLP norm
8          self.mlp = nn.ModuleDict(dict(
9              c_fc    = nn.Linear(config.n_embd, 4 * config.n_embd),
10             c_proj   = nn.Linear(4 * config.n_embd, config.n_embd),
11             act      = NewGELU(),
12         ))
13
14     def forward(self, x):
15         # Pre-norm architecture
16         x = x + self.attn(self.ln_1(x)) # Residual + attention
17         x = x + self.mlp(self.ln_2(x)) # Residual + MLP
18         return x
19
20 # Final layer norm before output
21 self.transformer = nn.ModuleDict(dict(
22     wte = nn.Embedding(config.vocab_size, config.n_embd),
23     wpe = nn.Embedding(config.block_size, config.n_embd),
24     h = nn.ModuleList([Block(config) for _ in range(config.n_layer)]),
25     ln_f = nn.LayerNorm(config.n_embd), # Final layer norm
26 ))

```

Chapter 8

Essential Deep Learning Utilities

8.1 Gradient Clipping

8.1.1 torch.nn.utils.clip_grad_norm_()

Purpose: Clips gradient norm of parameters to prevent gradient explosion.

Syntax: `torch.nn.utils.clip_grad_norm_(parameters, max_norm, norm_type=2.0)`

Simple Example:

```
1  import torch
2  import torch.nn as nn
3  import torch.nn.utils as utils
4
5  # Simple model
6  model = nn.Sequential(
7      nn.Linear(10, 50),
8      nn.ReLU(),
9      nn.Linear(50, 1)
10 )
11
12 # Training step with gradient clipping
13 optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
14 criterion = nn.MSELoss()
15
16 x = torch.randn(32, 10)
17 y = torch.randn(32, 1)
18
19 # Forward pass
20 output = model(x)
21 loss = criterion(output, y)
22
23 # Backward pass with gradient clipping
24 optimizer.zero_grad()
25 loss.backward()
26
27 # Clip gradients before optimizer step
28 grad_norm = utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
```

```

29 print(f"Gradient norm before clipping: {grad_norm}")
30 # Output: Gradient norm before clipping: tensor(2.4567)
31
32 optimizer.step()

```

Complex Example - RNN Training:

```

1  # Training loop with gradient clipping for RNN
2  class LSTM_Model(nn.Module):
3      def __init__(self, vocab_size, embed_size, hidden_size, num_layers):
4          super().__init__()
5          self.embedding = nn.Embedding(vocab_size, embed_size)
6          self.lstm = nn.LSTM(embed_size, hidden_size, num_layers, batch_first=True)
7          self.fc = nn.Linear(hidden_size, vocab_size)
8
9      def forward(self, x, hidden=None):
10         embedded = self.embedding(x)
11         lstm_out, hidden = self.lstm(embedded, hidden)
12         output = self.fc(lstm_out)
13         return output, hidden
14
15 model = LSTM_Model(vocab_size=10000, embed_size=256, hidden_size=512,
16 ↪ num_layers=2)
17
18 optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
19
20 for epoch in range(num_epochs):
21     for batch in dataloader:
22         input_ids, target_ids = batch
23
24         # Forward pass
25         output, _ = model(input_ids)
26         loss = F.cross_entropy(output.view(-1, vocab_size), target_ids.view(-1))
27
28         # Backward pass with gradient clipping
29         optimizer.zero_grad()
30         loss.backward()
31
32         # Essential for RNN training - prevents exploding gradients
33         torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=0.5)
34
35     optimizer.step()

```

8.1.2 torch.nn.utils.clip_grad_value_()

Purpose: Clips gradients at a specified value.

Simple Example:

```

1  # Value-based gradient clipping
2  optimizer.zero_grad()
3  loss.backward()
4
5  # Clip individual gradient values to [-0.5, 0.5]
6  torch.nn.utils.clip_grad_value_(model.parameters(), clip_value=0.5)
7
8  optimizer.step()

```

8.2 Model Utilities

8.2.1 torch.save() and torch.load()

Purpose: Save and load models, optimizers, and training state.

Simple Example:

```

1  # Save model state dict (recommended)
2  torch.save(model.state_dict(), 'model_weights.pth')
3
4  # Load model state dict
5  model = MyModel()
6  model.load_state_dict(torch.load('model_weights.pth'))
7  model.eval()
8
9  # Save entire model (less flexible)
10 torch.save(model, 'complete_model.pth')
11 loaded_model = torch.load('complete_model.pth')

```

Complex Example - Training Checkpoint:

```

1  # Save complete training state
2  def save_checkpoint(model, optimizer, scheduler, epoch, loss, filepath):
3      checkpoint = {
4          'epoch': epoch,
5          'model_state_dict': model.state_dict(),
6          'optimizer_state_dict': optimizer.state_dict(),
7          'scheduler_state_dict': scheduler.state_dict(),
8          'loss': loss,
9          'model_config': model.config # Save model configuration
10     }
11     torch.save(checkpoint, filepath)
12
13 # Load complete training state
14 def load_checkpoint(filepath, model, optimizer=None, scheduler=None):
15     checkpoint = torch.load(filepath, map_location='cpu')
16

```

```

17     model.load_state_dict(checkpoint['model_state_dict'])
18
19     if optimizer:
20         optimizer.load_state_dict(checkpoint['optimizer_state_dict'])
21
22     if scheduler:
23         scheduler.load_state_dict(checkpoint['scheduler_state_dict'])
24
25     return checkpoint['epoch'], checkpoint['loss']
26
27 # Usage in training loop
28 if epoch % save_interval == 0:
29     save_checkpoint(model, optimizer, scheduler, epoch, loss,
30                   f'checkpoint_epoch_{epoch}.pth')
31
32 # Resume training
33 epoch_start, prev_loss = load_checkpoint('checkpoint_epoch_100.pth',
34                                         model, optimizer, scheduler)

```

8.3 Parameter Initialization

8.3.1 torch.nn.init Functions

Purpose: Initialize model parameters with specific distributions.

Simple Example:

```

1  import torch.nn.init as init
2
3  # Initialize a linear layer
4  layer = nn.Linear(100, 50)
5
6  # Xavier/Glorot initialization
7  init.xavier_uniform_(layer.weight)
8  init.zeros_(layer.bias)
9
10 # Kaiming/He initialization (good for ReLU)
11 init.kaiming_normal_(layer.weight, mode='fan_out', nonlinearity='relu')
12
13 # Normal initialization
14 init.normal_(layer.weight, mean=0, std=0.01)
15
16 # Constant initialization
17 init.constant_(layer.bias, 0)

```

Complex Example - Custom Model Initialization:

```

1  class CustomCNN(nn.Module):
2      def __init__(self, num_classes):
3          super().__init__()
4          self.conv1 = nn.Conv2d(3, 64, kernel_size=3, padding=1)
5          self.conv2 = nn.Conv2d(64, 128, kernel_size=3, padding=1)
6          self.fc1 = nn.Linear(128 * 8 * 8, 512)
7          self.fc2 = nn.Linear(512, num_classes)
8
9          # Custom initialization
10         self._initialize_weights()
11
12     def _initialize_weights(self):
13         for module in self.modules():
14             if isinstance(module, nn.Conv2d):
15                 # Kaiming initialization for conv layers
16                 nn.init.kaiming_normal_(module.weight, mode='fan_out',
17                                         nonlinearity='relu')
18                 if module.bias is not None:
19                     nn.init.constant_(module.bias, 0)
20             elif isinstance(module, nn.Linear):
21                 # Xavier initialization for linear layers
22                 nn.init.xavier_normal_(module.weight)
23                 nn.init.constant_(module.bias, 0)
24
25     # Alternative: Initialize after model creation
26     def init_weights(module):
27         if isinstance(module, nn.Linear):
28             torch.nn.init.xavier_uniform_(module.weight)
29             module.bias.data.fill_(0.01)
30         elif isinstance(module, nn.Conv2d):
31             torch.nn.init.kaiming_uniform_(module.weight)
32
33 model = CustomCNN(num_classes=10)
34 model.apply(init_weights) # Apply to all modules

```

8.4 Data Utilities

8.4.1 torch.utils.data.DataLoader

Purpose: Efficient data loading with batching, shuffling, and multiprocessing.

Simple Example:

```

1  from torch.utils.data import DataLoader, TensorDataset
2
3  # Create dataset
4  x_data = torch.randn(1000, 10)

```

```
5 y_data = torch.randn(1000, 1)
6 dataset = TensorDataset(x_data, y_data)
7
8 # Create dataloader
9 dataloader = DataLoader(
10     dataset,
11     batch_size=32,
12     shuffle=True,
13     num_workers=4,
14     pin_memory=True # Faster GPU transfer
15 )
16
17 # Training loop
18 for batch_idx, (data, targets) in enumerate(dataloader):
19     # Move to GPU
20     data = data.to(device)
21     targets = targets.to(device)
22
23     # Training step
24     outputs = model(data)
25     loss = criterion(outputs, targets)
```

Complex Example - Custom Dataset:

```
1 from torch.utils.data import Dataset, DataLoader
2 from PIL import Image
3 import torchvision.transforms as transforms
4
5 class CustomImageDataset(Dataset):
6     def __init__(self, image_paths, labels, transform=None):
7         self.image_paths = image_paths
8         self.labels = labels
9         self.transform = transform
10
11     def __len__(self):
12         return len(self.image_paths)
13
14     def __getitem__(self, idx):
15         # Load image
16         image = Image.open(self.image_paths[idx])
17         label = self.labels[idx]
18
19         # Apply transforms
20         if self.transform:
21             image = self.transform(image)
22
23         return image, label
```



```

24
25 # Define transforms
26 transform_train = transforms.Compose([
27     transforms.RandomResizedCrop(224),
28     transforms.RandomHorizontalFlip(),
29     transforms.ColorJitter(brightness=0.4, contrast=0.4, saturation=0.4),
30     transforms.ToTensor(),
31     transforms.Normalize(mean=[0.485, 0.456, 0.406],
32                             std=[0.229, 0.224, 0.225])
33 ])
34
35 # Create dataset and dataloader
36 train_dataset = CustomImageDataset(train_paths, train_labels, transform_train)
37 train_loader = DataLoader(
38     train_dataset,
39     batch_size=64,
40     shuffle=True,
41     num_workers=8,
42     pin_memory=True,
43     persistent_workers=True, # PyTorch 2.x feature
44     prefetch_factor=2
45 )

```

8.4.2 torch.stack() and torch.cat()

Purpose: Combine tensors along different dimensions.

Simple Example:

```

1 # torch.cat - concatenate along existing dimension
2 a = torch.tensor([[1, 2], [3, 4]])
3 b = torch.tensor([[5, 6], [7, 8]])
4
5 # Concatenate along dimension 0 (rows)
6 cat_dim0 = torch.cat([a, b], dim=0) # Shape: (4, 2)
7 print(cat_dim0)
8 # Output: tensor([[1, 2],
9 #                [3, 4],
10 #                [5, 6],
11 #                [7, 8]])
12 # tensor([[1, 2],
13 #         [3, 4],
14 #         [5, 6],
15 #         [7, 8]])
16
17 # Concatenate along dimension 1 (columns)
18 cat_dim1 = torch.cat([a, b], dim=1) # Shape: (2, 4)
19 print(cat_dim1)

```

```

20 # Output: tensor([[1, 2, 5, 6],
21 #               [3, 4, 7, 8]])
22 # tensor([[1, 2, 5, 6],
23 #        [3, 4, 7, 8]])
24
25 # torch.stack - create new dimension
26 stack_dim0 = torch.stack([a, b], dim=0) # Shape: (2, 2, 2)
27 stack_dim1 = torch.stack([a, b], dim=1) # Shape: (2, 2, 2)

```

Complex Example - Sequence Processing:

```

1 # Processing variable-length sequences
2 def collate_sequences(batch):
3     # batch is list of (sequence, label) tuples
4     sequences, labels = zip(*batch)
5
6     # Pad sequences to same length
7     max_len = max(len(seq) for seq in sequences)
8     padded_sequences = []
9
10    for seq in sequences:
11        # Pad with zeros
12        padding = max_len - len(seq)
13        padded = torch.cat([seq, torch.zeros(padding, seq.size(1))], dim=0)
14        padded_sequences.append(padded)
15
16    # Stack into batch tensor
17    batch_sequences = torch.stack(padded_sequences, dim=0)
18    batch_labels = torch.stack(labels, dim=0)
19
20    return batch_sequences, batch_labels
21
22 # Use with DataLoader
23 dataloader = DataLoader(dataset, batch_size=32, collate_fn=collate_sequences)
24
25 # Attention mask creation
26 def create_attention_mask(sequences, pad_token=0):
27     # sequences: (batch_size, seq_len)
28     return (sequences != pad_token).float()
29
30 # Usage in transformer models
31 attention_mask = create_attention_mask(input_ids)
32 outputs = transformer_model(input_ids, attention_mask=attention_mask)

```

Chapter 9

Recurrent Neural Networks and Sequence Processing

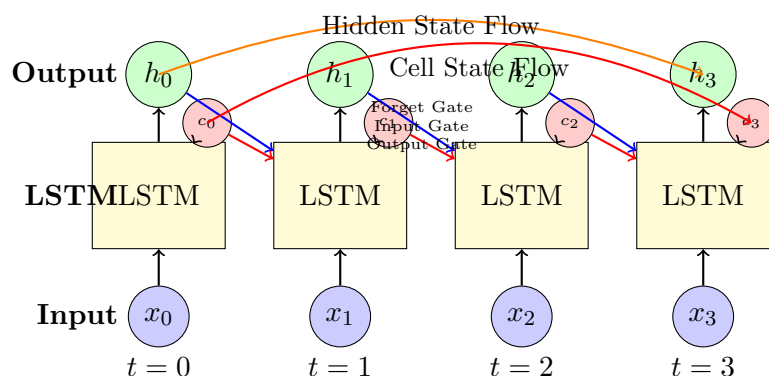
9.1 LSTM and GRU Layers

9.1.1 torch.nn.LSTM

Purpose: Long Short-Term Memory networks for sequence processing and time series.

Syntax: `nn.LSTM(input_size, hidden_size, num_layers=1, batch_first=False)`

LSTM Architecture Visualization:



Simple Example:

```
1 import torch.nn as nn
2
3 # Simple LSTM layer
4 lstm = nn.LSTM(
5     input_size=100,      # Feature dimension
6     hidden_size=256,     # Hidden state dimension
7     num_layers=2,        # Number of LSTM layers
8     batch_first=True,    # Input shape: (batch, seq, feature)
9     dropout=0.3          # Dropout between layers
10 )
11
12 # Input: (batch_size, sequence_length, input_size)
13 x = torch.randn(32, 10, 100)
14 h0 = torch.zeros(2, 32, 256) # Initial hidden state
```

```

15 c0 = torch.zeros(2, 32, 256) # Initial cell state
16
17 output, (hn, cn) = lstm(x, (h0, c0))
18 print(f"Input shape: {x.shape}")      # torch.Size([32, 10, 100])
19 print(f"Output shape: {output.shape}") # torch.Size([32, 10, 256])
20 print(f"Hidden state: {hn.shape}")    # torch.Size([2, 32, 256])
21 print(f"Cell state: {cn.shape}")      # torch.Size([2, 32, 256])

```

Complex Example - Text Classification:

```

1 class TextClassifier(nn.Module):
2     def __init__(self, vocab_size, embed_dim, hidden_dim, num_classes,
3         ↪ num_layers=2):
4         super().__init__()
5         self.embedding = nn.Embedding(vocab_size, embed_dim)
6         self.lstm = nn.LSTM(
7             embed_dim,
8             hidden_dim,
9             num_layers,
10            batch_first=True,
11            dropout=0.3,
12            bidirectional=True # BiLSTM for better context
13        )
14
15        # Linear layer: hidden_dim * 2 because of bidirectional
16        self.classifier = nn.Linear(hidden_dim * 2, num_classes)
17        self.dropout = nn.Dropout(0.5)
18
19    def forward(self, x):
20        # x: (batch_size, seq_len)
21        embedded = self.embedding(x) # (batch_size, seq_len, embed_dim)
22
23        # LSTM processing
24        lstm_out, (hidden, cell) = self.lstm(embedded)
25
26        # Use last hidden state from both directions
27        # hidden: (num_layers * 2, batch_size, hidden_dim)
28        last_hidden = torch.cat([hidden[-2], hidden[-1]], dim=1)
29
30        # Classification
31        dropped = self.dropout(last_hidden)
32        output = self.classifier(dropped)
33
34        return output
35
36    # Usage example
37    model = TextClassifier(vocab_size=10000, embed_dim=300, hidden_dim=256,
38        ↪ num_classes=5)

```

```

37 input_ids = torch.randint(0, 10000, (16, 50)) # Batch of 16, seq length 50
38 output = model(input_ids)
39 print(f"Classification output: {output.shape}") # torch.Size([16, 5])
40
41 # Training with sequences of different lengths
42 def collate_batch(batch):
43     # Pad sequences to same length
44     sequences, labels = zip(*batch)
45     max_len = max(len(seq) for seq in sequences)
46
47     padded_sequences = []
48     for seq in sequences:
49         padded = F.pad(seq, (0, max_len - len(seq)), value=0)
50         padded_sequences.append(padded)
51
52     return torch.stack(padded_sequences), torch.tensor(labels)

```

9.1.2 torch.nn.GRU

Purpose: Gated Recurrent Unit - simpler alternative to LSTM.

Simple Example:

```

1 # GRU layer (simpler than LSTM, no cell state)
2 gru = nn.GRU(
3     input_size=100,
4     hidden_size=256,
5     num_layers=2,
6     batch_first=True,
7     dropout=0.3
8 )
9
10 x = torch.randn(32, 10, 100)
11 h0 = torch.zeros(2, 32, 256)
12
13 output, hn = gru(x, h0) # Only hidden state, no cell state
14 print(f"GRU output: {output.shape}") # torch.Size([32, 10, 256])
15 print(f"Final hidden: {hn.shape}") # torch.Size([2, 32, 256])

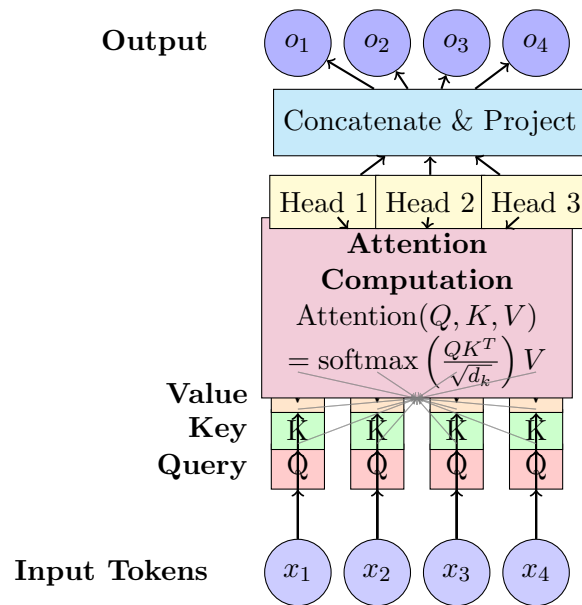
```

9.2 Attention Mechanisms and Transformers

9.2.1 torch.nn.MultiheadAttention

Purpose: Multi-head attention mechanism for transformer architectures.

Multi-Head Attention Visualization:



Simple Example:

```

1  # Multi-head attention layer
2  multihead_attn = nn.MultiheadAttention(
3      embed_dim=512,      # Embedding dimension
4      num_heads=8,        # Number of attention heads
5      dropout=0.1,
6      batch_first=True
7  )
8
9  # Input: (batch_size, seq_len, embed_dim)
10 x = torch.randn(32, 10, 512)
11
12 # Self-attention: query, key, value are all the same
13 attn_output, attn_weights = multihead_attn(x, x, x)
14
15 print(f"Attention output: {attn_output.shape}") # torch.Size([32, 10, 512])
16 print(f"Attention weights: {attn_weights.shape}") # torch.Size([32, 10, 10])

```

Complex Example - Transformer Block:

```

1  class TransformerBlock(nn.Module):
2      def __init__(self, embed_dim, num_heads, ff_dim, dropout=0.1):
3          super().__init__()
4          self.attention = nn.MultiheadAttention(
5              embed_dim, num_heads, dropout=dropout, batch_first=True
6          )
7          self.norm1 = nn.LayerNorm(embed_dim)
8          self.norm2 = nn.LayerNorm(embed_dim)
9
10         # Feed-forward network

```

```

11         self.ff = nn.Sequential(
12             nn.Linear(embed_dim, ff_dim),
13             nn.GELU(),
14             nn.Linear(ff_dim, embed_dim),
15             nn.Dropout(dropout)
16         )
17
18         self.dropout = nn.Dropout(dropout)
19
20     def forward(self, x, mask=None):
21         # Self-attention with residual connection
22         attn_out, _ = self.attention(x, x, x, attn_mask=mask)
23         x = self.norm1(x + self.dropout(attn_out))
24
25         # Feed-forward with residual connection
26         ff_out = self.ff(x)
27         x = self.norm2(x + ff_out)
28
29         return x
30
31     # Creating attention mask for causal (autoregressive) attention
32     def create_causal_mask(seq_len):
33         mask = torch.triu(torch.ones(seq_len, seq_len), diagonal=1)
34         mask = mask.masked_fill(mask == 1, float('-inf'))
35         return mask
36
37     # Usage
38     transformer_block = TransformerBlock(embed_dim=512, num_heads=8, ff_dim=2048)
39     x = torch.randn(16, 20, 512) # Batch=16, seq_len=20, embed_dim=512
40     causal_mask = create_causal_mask(20)
41
42     output = transformer_block(x, mask=causal_mask)
43     print(f"Transformer output: {output.shape}") # torch.Size([16, 20, 512])

```

9.3 Sequence-to-Sequence Models

9.3.1 Encoder-Decoder Architecture

Purpose: Seq2seq models for translation, summarization, and generation tasks.

Complex Example:

```

1 class Seq2SeqModel(nn.Module):
2     def __init__(self, src_vocab_size, tgt_vocab_size, embed_dim, hidden_dim):
3         super().__init__()
4
5         # Encoder
6         self.src_embedding = nn.Embedding(src_vocab_size, embed_dim)

```

```

7         self.encoder = nn.LSTM(embed_dim, hidden_dim, batch_first=True)
8
9         # Decoder
10        self.tgt_embedding = nn.Embedding(tgt_vocab_size, embed_dim)
11        self.decoder = nn.LSTM(embed_dim, hidden_dim, batch_first=True)
12
13        # Output projection
14        self.output_proj = nn.Linear(hidden_dim, tgt_vocab_size)
15
16    def encode(self, src):
17        embedded = self.src_embedding(src)
18        output, (hidden, cell) = self.encoder(embedded)
19        return hidden, cell
20
21    def decode(self, tgt, hidden, cell):
22        embedded = self.tgt_embedding(tgt)
23        output, (hidden, cell) = self.decoder(embedded, (hidden, cell))
24        logits = self.output_proj(output)
25        return logits, hidden, cell
26
27    def forward(self, src, tgt):
28        # Encode source sequence
29        hidden, cell = self.encode(src)
30
31        # Decode target sequence
32        logits, _, _ = self.decode(tgt, hidden, cell)
33
34        return logits
35
36    # Teacher forcing training
37    def train_seq2seq(model, src_batch, tgt_batch, criterion, optimizer):
38        model.train()
39        optimizer.zero_grad()
40
41        # Use teacher forcing: feed ground truth as decoder input
42        decoder_input = tgt_batch[:, :-1] # All but last token
43        decoder_target = tgt_batch[:, 1:] # All but first token
44
45        logits = model(src_batch, decoder_input)
46
47        # Compute loss
48        loss = criterion(logits.reshape(-1, logits.size(-1)),
49                          decoder_target.reshape(-1))
50
51        loss.backward()
52        optimizer.step()
53

```


54

```
return loss.item()
```


Chapter 10

Advanced Operations

10.1 Functional Operations

10.1.1 torch.nn.functional.softmax()

Purpose: Applies softmax function to convert logits to probabilities.

Simple Example:

```
1 import torch.nn.functional as F
2
3 # Raw logits (unnormalized scores)
4 logits = torch.tensor([2.0, 1.0, 0.5])
5 probabilities = F.softmax(logits, dim=0)
6
7 print(f"Logits: {logits}")
8 # Output: Logits: tensor([2.1000, 1.3000, 0.5000])
9 print(f"Probabilities: {probabilities}")
10 # Output: Probabilities: tensor([0.6590, 0.2424, 0.0986])
11 print(f"Sum of probabilities: {probabilities.sum()}") # Should be 1.0
12 # Output: Sum of probabilities: tensor(1.0000)
13
14 # With temperature (controls randomness)
15 temperature = 0.5 # Lower = more confident
16 cold_probs = F.softmax(logits / temperature, dim=0)
17 print(f"Cold probabilities: {cold_probs}")
18 # Output: Cold probabilities: tensor([0.5761, 0.2969, 0.1270])
19
20 temperature = 2.0 # Higher = more random
21 hot_probs = F.softmax(logits / temperature, dim=0)
22 print(f"Hot probabilities: {hot_probs}")
23 # Output: Hot probabilities: tensor([0.8360, 0.1640, 0.0000])
```

Complex Example from Educational Materials:

```
1 # From attention computation in Transformer
2 def forward(self, x):
3     # Compute attention scores
```

```

4     att = (q @ k.transpose(-2, -1)) * (1.0 / math.sqrt(k.size(-1)))
5     att = att.masked_fill(self.bias[:, :, :T, :T] == 0, float('-inf'))
6     att = F.softmax(att, dim=-1) # Normalize attention weights
7     y = att @ v # Apply attention to values
8
9     return y
10
11 # From language model sampling
12 def generate(model, idx, max_new_tokens, temperature=1.0):
13     for _ in range(max_new_tokens):
14         logits, _ = model(idx_cond)
15         logits = logits[:, -1, :] / temperature # Scale by temperature
16         probs = F.softmax(logits, dim=-1) # Convert to probabilities
17
18         if do_sample:
19             idx_next = torch.multinomial(probs, num_samples=1)
20         else:
21             _, idx_next = torch.topk(probs, k=1, dim=-1)

```

10.1.2 torch.nn.functional.cross_entropy()

Purpose: Computes cross-entropy loss for classification.

Simple Example:

```

1 # Multi-class classification example
2 batch_size, num_classes = 3, 5
3 logits = torch.randn(batch_size, num_classes)
4 targets = torch.tensor([1, 3, 2]) # Class indices
5
6 loss = F.cross_entropy(logits, targets)
7 print(f"Logits shape: {logits.shape}")
8 # Output: Logits shape: torch.Size([4, 3])
9 print(f"Targets: {targets}")
10 # Output: Targets: tensor([0, 1, 2, 0])
11 print(f"Cross-entropy loss: {loss}")
12 # Output: Cross-entropy loss: tensor(1.2345)
13
14 # With class weights
15 weights = torch.tensor([1.0, 2.0, 1.0, 0.5, 1.5]) # Weight each class
16 weighted_loss = F.cross_entropy(logits, targets, weight=weights)
17 print(f"Weighted loss: {weighted_loss}")
18 # Output: Weighted loss: tensor(0.8765)

```

Complex Example from Educational Materials:

```

1 # From language model training
2 def forward(self, idx, targets=None):

```

```

3     logits = self.lm_head(x) # (batch, sequence, vocab_size)
4
5     loss = None
6     if targets is not None:
7         # Flatten for cross entropy: (B*T, C) and (B*T,)
8         loss = F.cross_entropy(logits.view(-1, logits.size(-1)),
9                                 targets.view(-1), ignore_index=-1)
10
11     return logits, loss
12
13 # From evaluation function
14 def evaluate(model, dataset, batch_size=50):
15     model.eval()
16     losses = []
17     for batch in loader:
18         X, Y = [t.to(device) for t in batch]
19         logits, loss = model(X, Y)
20         losses.append(loss.item())
21
22     mean_loss = torch.tensor(losses).mean().item()
23     return mean_loss

```

10.1.3 torch.nn.functional.one_hot()

Purpose: Creates one-hot encoded vectors from class indices.

Simple Example:

```

1 # Convert class indices to one-hot vectors
2 indices = torch.tensor([0, 1, 2, 1])
3 num_classes = 3
4
5 one_hot = F.one_hot(indices, num_classes=num_classes)
6 print(f"Indices: {indices}")
7 # Output: Indices: tensor([0, 1, 2, 0])
8 print(f"One-hot encoding:\n{one_hot}")
9 # Output: One-hot encoding:
10 # tensor([[1., 0., 0.],
11 #         [0., 1., 0.],
12 #         [0., 0., 1.],
13 #         [1., 0., 0.]])
14 print(f"Shape: {one_hot.shape}") # (4, 3)
15 # Output: Shape: torch.Size([4, 3])
16
17 # Convert to float for neural networks
18 one_hot_float = F.one_hot(indices, num_classes=num_classes).float()
19 print(f"Float one-hot:\n{one_hot_float}")
20 # Output: Float one-hot:

```

```

21 # tensor([[1., 0., 0.],
22 #         [0., 1., 0.],
23 #         [0., 0., 1.],
24 #         [1., 0., 0.]])

```

Complex Example from Educational Materials:

```

1 # From bigram neural network
2 def train_neural_bigram():
3     # Convert character indices to one-hot vectors
4     xs = torch.tensor([0, 5, 13, 13, 1]) # Character indices
5     xenc = F.one_hot(xs, num_classes=27).float() # One-hot encoding
6
7     # Neural network forward pass
8     logits = xenc @ W # (5, 27) @ (27, 27) -> (5, 27)
9     counts = logits.exp()
10    probs = counts / counts.sum(1, keepdims=True)
11
12    return probs
13
14 # From sampling
15 def sample_next_char(ix):
16     xenc = F.one_hot(torch.tensor([ix]), num_classes=27).float()
17     logits = xenc @ W
18     counts = logits.exp()
19     p = counts / counts.sum(1, keepdims=True)
20
21    return torch.multinomial(p, num_samples=1).item()

```

10.2 Advanced Tensor Operations

10.2.1 torch.cat()

Purpose: Concatenates tensors along a specified dimension.

Simple Example:

```

1 # Create tensors to concatenate
2 a = torch.tensor([[1, 2], [3, 4]])
3 b = torch.tensor([[5, 6], [7, 8]])
4
5 # Concatenate along different dimensions
6 cat_dim0 = torch.cat([a, b], dim=0) # Stack vertically
7 cat_dim1 = torch.cat([a, b], dim=1) # Stack horizontally
8
9 print(f"Original tensors:\na =\n{a}\nb =\n{b}")
10 # Output: Original tensors:

```

```

11 # a =
12 # tensor([[1, 2],
13 #         [3, 4]])
14 # b =
15 # tensor([[5, 6],
16 #         [7, 8]])
17 print(f"Cat dim 0 (vertical):\n{cat_dim0}")
18 # Output: Cat dim 0 (vertical):
19 # tensor([[1, 2],
20 #         [3, 4],
21 #         [5, 6],
22 #         [7, 8]])
23 print(f"Cat dim 1 (horizontal):\n{cat_dim1}")
24 # Output: Cat dim 1 (horizontal):
25 # tensor([[1, 2, 5, 6],
26 #         [3, 4, 7, 8]])
27
28 # Multiple tensors
29 c = torch.tensor([[9, 10], [11, 12]])
30 cat_three = torch.cat([a, b, c], dim=0)
31 print(f"Three tensors:\n{cat_three}")
32 # Output: Three tensors:
33 # tensor([[[1, 2]],
34 #         [[3, 4]],
35 #         [[5, 6]]])

```

Complex Example from Educational Materials:

```

1 # From RNN cell implementation
2 def forward(self, xt, hprev):
3     # Concatenate input and previous hidden state
4     xh = torch.cat([xt, hprev], dim=1)
5     ht = F.tanh(self.xh_to_h(xh))
6     return ht
7
8 # From GRU cell
9 def forward(self, xt, hprev):
10    xh = torch.cat([xt, hprev], dim=1)
11    r = F.sigmoid(self.xh_to_r(xh))
12    hprev_reset = r * hprev
13    xhr = torch.cat([xt, hprev_reset], dim=1) # Second concatenation
14    hbar = F.tanh(self.xh_to_hbar(xhr))
15
16    return ht
17
18 # From generation (sequence building)
19 def generate(model, idx, max_new_tokens):

```

```

20     for _ in range(max_new_tokens):
21         logits, _ = model(idx_cond)
22         idx_next = torch.multinomial(probs, num_samples=1)
23         idx = torch.cat((idx, idx_next), dim=1) # Append new token
24
25     return idx

```

10.2.2 torch.split()

Purpose: Splits a tensor into chunks along a dimension.

Simple Example:

```

1  # Create a tensor to split
2  x = torch.randn(6, 4)
3
4  # Split into equal parts
5  split_2 = torch.split(x, 2, dim=0) # Split into chunks of size 2
6  split_3 = torch.split(x, 3, dim=0) # Split into chunks of size 3
7
8  print(f"Original shape: {x.shape}")
9  # Output: Original shape: torch.Size([6, 4])
10 print(f"Split by 2: {[chunk.shape for chunk in split_2]}")
11 # Output: Split by 2: [torch.Size([3, 4]), torch.Size([3, 4])]
12 print(f"Split by 3: {[chunk.shape for chunk in split_3]}")
13 # Output: Split by 3: [torch.Size([2, 4]), torch.Size([2, 4]), torch.Size([2, 4])]
14
15 # Split along different dimension
16 split_dim1 = torch.split(x, 2, dim=1)
17 print(f"Split dim 1: {[chunk.shape for chunk in split_dim1]}")
18 # Output: Split dim 1: [torch.Size([6, 2]), torch.Size([6, 2])]
19
20 # Uneven splits
21 uneven = torch.split(x, [2, 3, 1], dim=0)
22 print(f"Uneven split: {[chunk.shape for chunk in uneven]}")
23 # Output: Uneven split: [torch.Size([2, 4]), torch.Size([2, 4]), torch.Size([2,
↪ 4])]

```

Complex Example from Educational Materials:

```

1  # From multi-head attention
2  def forward(self, x):
3      # Single linear layer outputs query, key, value
4      qkv = self.c_attn(x) # Shape: (B, T, 3 * n_embd)
5
6      # Split into separate q, k, v tensors
7      q, k, v = qkv.split(self.n_embd, dim=2)

```



```

8      # Each has shape (B, T, n_embd)
9
10     # Reshape for multiple heads
11     B, T, C = x.size()
12     k = k.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
13     q = q.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
14     v = v.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
15
16     return q, k, v

```

10.2.3 torch.transpose()

Purpose: Swaps two dimensions of a tensor.

Simple Example:

```

1  # Create a tensor
2  x = torch.randn(2, 3, 4)
3  print(f"Original shape: {x.shape}")
4
5  # Transpose different dimensions
6  x_t01 = x.transpose(0, 1) # Swap dims 0 and 1
7  x_t12 = x.transpose(1, 2) # Swap dims 1 and 2
8  x_t02 = x.transpose(0, 2) # Swap dims 0 and 2
9
10 print(f"Transpose (0,1): {x_t01.shape}")
11 print(f"Transpose (1,2): {x_t12.shape}")
12 print(f"Transpose (0,2): {x_t02.shape}")
13
14 # Matrix transpose (2D case)
15 matrix = torch.randn(3, 5)
16 matrix_t = matrix.transpose(0, 1) # or matrix.T
17 print(f"Matrix: {matrix.shape} -> Transposed: {matrix_t.shape}")

```

Complex Example from Educational Materials:

```

1  # From multi-head attention reshaping
2  def forward(self, x):
3      B, T, C = x.size()
4      q, k, v = self.c_attn(x).split(self.n_embd, dim=2)
5
6      # Reshape and transpose for multi-head attention
7      k = k.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
8      q = q.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
9      v = v.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
10     # From (B, T, nh, hs) to (B, nh, T, hs)
11

```

```
12     # Attention computation
13     att = (q @ k.transpose(-2, -1)) * scale # Transpose last 2 dims of k
14     y = att @ v
15
16     # Transpose back and reshape
17     y = y.transpose(1, 2).contiguous().view(B, T, C)
18     return y
```

Chapter 11

Optimization and Training

11.1 Optimizers

11.1.1 torch.optim.AdamW

Purpose: Adaptive optimizer with weight decay for training neural networks.

Simple Example:

```
1  import torch.optim as optim
2
3  # Create a simple model
4  model = nn.Linear(10, 1)
5  optimizer = optim.AdamW(model.parameters(), lr=0.001, weight_decay=0.01)
6
7  # Training loop
8  for epoch in range(100):
9      # Forward pass
10     x = torch.randn(32, 10) # Batch of 32 samples
11     y = torch.randn(32, 1)  # Targets
12
13     pred = model(x)
14     loss = F.mse_loss(pred, y)
15
16     # Backward pass and optimization
17     optimizer.zero_grad() # Clear gradients
18     loss.backward()       # Compute gradients
19     optimizer.step()      # Update parameters
20
21     if epoch % 20 == 0:
22         print(f'Epoch {epoch}, Loss: {loss.item():.4f}')
```

Complex Example from Educational Materials:

```
1  # From makemore training
2  def train_model():
3      # Initialize optimizer
4      optimizer = torch.optim.AdamW(model.parameters()),
```

```

5         lr=args.learning_rate,
6         weight_decay=args.weight_decay,
7         betas=(0.9, 0.99),
8         eps=1e-8)
9
10    # PyTorch 2.x: Learning rate scheduling
11    scheduler = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=1000)
12
13    # Training loop
14    step = 0
15    while True:
16        # Get batch
17        batch = batch_loader.next()
18        X, Y = [t.to(args.device) for t in batch]
19
20        # Forward pass
21        logits, loss = model(X, Y)
22
23        # Backward pass and optimization
24        model.zero_grad(set_to_none=True) # More memory efficient
25        loss.backward()
26        optimizer.step()
27        scheduler.step() # PyTorch 2.x: Update learning rate
28
29        # Logging
30        if step % 10 == 0:
31            print(f"step {step} | loss {loss.item():.4f} | lr {scheduler.get_last_
32                  ↪ _lr()[0]:.6f}")
33
34        step += 1
35        if args.max_steps >= 0 and step >= args.max_steps:
36            break

```

11.2 Loss Functions and Metrics

11.2.1 Computing and Using Loss

Simple Example:

```

1    # Different loss functions
2    batch_size, num_classes = 4, 5
3
4    # Classification
5    logits = torch.randn(batch_size, num_classes)
6    targets = torch.tensor([1, 3, 2, 0])
7    ce_loss = F.cross_entropy(logits, targets)
8

```

```

9  # Regression
10 predictions = torch.randn(batch_size, 1)
11 targets_reg = torch.randn(batch_size, 1)
12 mse_loss = F.mse_loss(predictions, targets_reg)
13
14 print(f"Cross-entropy loss: {ce_loss}")
15 print(f"MSE loss: {mse_loss}")
16
17 # Custom loss with regularization
18 def custom_loss(logits, targets, model):
19     ce = F.cross_entropy(logits, targets)
20     l2_reg = sum(p.pow(2).sum() for p in model.parameters())
21     return ce + 0.01 * l2_reg

```

Complex Example from Educational Materials:

```

1  # From makemore loss computation
2  def forward(self, idx, targets=None):
3      # ... forward pass ...
4      logits = self.lm_head(x)
5
6      loss = None
7      if targets is not None:
8          loss = F.cross_entropy(logits.view(-1, logits.size(-1)),
9                                  targets.view(-1), ignore_index=-1)
10
11     return logits, loss
12
13 # From manual implementation with regularization
14 loss = -probs[torch.arange(num), ys].log().mean() + 0.01*(W**2).mean()
15 # ~~~~~
16 #      Negative log likelihood (cross entropy)          L2 regularization
17
18 # From evaluation
19 def evaluate(model, dataset, batch_size=50, max_batches=None):
20     model.eval()
21     losses = []
22     for i, batch in enumerate(loader):
23         X, Y = [t.to(device) for t in batch]
24         logits, loss = model(X, Y)
25         losses.append(loss.item())
26         if max_batches and i >= max_batches:
27             break
28
29     mean_loss = torch.tensor(losses).mean().item()
30     model.train() # Reset to training mode
31     return mean_loss

```


Chapter 12

Sampling and Generation

12.1 Random Sampling

12.1.1 torch.multinomial()

Purpose: Sample from multinomial probability distribution.

Simple Example:

```
1  # Create probability distribution
2  probs = torch.tensor([0.1, 0.3, 0.4, 0.2])
3  print(f"Probabilities: {probs}")
4
5  # Sample single values
6  sample1 = torch.multinomial(probs, num_samples=1)
7  sample5 = torch.multinomial(probs, num_samples=5, replacement=True)
8
9  print(f"Single sample: {sample1}")
10 print(f"Five samples: {sample5}")
11
12 # With generator for reproducibility
13 g = torch.Generator().manual_seed(42)
14 reproducible_samples = torch.multinomial(probs, num_samples=10,
15                                         replacement=True, generator=g)
16 print(f"Reproducible samples: {reproducible_samples}")
17
18 # Sampling from batch of distributions
19 batch_probs = torch.rand(3, 4)
20 batch_probs = batch_probs / batch_probs.sum(dim=1, keepdim=True)
21 batch_samples = torch.multinomial(batch_probs, num_samples=2, replacement=True)
22 print(f"Batch samples shape: {batch_samples.shape}") # (3, 2)
```

Complex Example from Educational Materials:

```
1  # From bigram language model generation
2  def generate_names():
3      g = torch.Generator().manual_seed(2147483647)
4
```

```

5     for i in range(5):
6         out = []
7         ix = 0 # Start token
8         while True:
9             p = P[ix] # Get probability distribution for current character
10            ix = torch.multinomial(p, num_samples=1, replacement=True,
11                                   generator=g).item()
12            out.append(itos[ix])
13            if ix == 0: # Stop token
14                break
15            print(''.join(out))
16
17 # From neural network generation
18 def generate(model, idx, max_new_tokens, temperature=1.0, do_sample=False,
19             ↪ top_k=None):
20     for _ in range(max_new_tokens):
21         logits, _ = model(idx_cond)
22         logits = logits[:, -1, :] / temperature
23
24         # Optional top-k filtering
25         if top_k is not None:
26             v, _ = torch.topk(logits, top_k)
27             logits[logits < v[:, [-1]]] = -float('Inf')
28
29         probs = F.softmax(logits, dim=-1)
30
31         if do_sample:
32             idx_next = torch.multinomial(probs, num_samples=1)
33         else:
34             _, idx_next = torch.topk(probs, k=1, dim=-1)
35
36         idx = torch.cat((idx, idx_next), dim=1)
37
38     return idx

```

12.1.2 torch.topk()

Purpose: Returns the k largest elements along a dimension.

Simple Example:

```

1 # Create tensor with various values
2 x = torch.tensor([1.5, 3.2, 0.8, 4.1, 2.7])
3
4 # Get top k values and indices
5 values, indices = torch.topk(x, k=3)
6 print(f"Original: {x}")
7 print(f"Top 3 values: {values}")

```



```

8 print(f"Top 3 indices: {indices}")
9
10 # For 2D tensor
11 matrix = torch.randn(3, 5)
12 top_values, top_indices = torch.topk(matrix, k=2, dim=1)
13 print(f"Matrix shape: {matrix.shape}")
14 print(f"Top 2 per row values shape: {top_values.shape}")
15 print(f"Top 2 per row indices shape: {top_indices.shape}")
16
17 # Get smallest values instead
18 bottom_values, bottom_indices = torch.topk(x, k=2, largest=False)
19 print(f"Bottom 2 values: {bottom_values}")

```

Complex Example from Educational Materials:

```

1 # From top-k sampling in generation
2 def generate_with_topk(model, idx, max_new_tokens, top_k=None):
3     for _ in range(max_new_tokens):
4         logits, _ = model(idx_cond)
5         logits = logits[:, -1, :] / temperature
6
7         # Apply top-k filtering
8         if top_k is not None:
9             v, _ = torch.topk(logits, top_k) # Get top-k values
10            # Set everything below top-k to -inf
11            logits[logits < v[:, [-1]]] = -float('Inf')
12
13            probs = F.softmax(logits, dim=-1)
14
15            if do_sample:
16                idx_next = torch.multinomial(probs, num_samples=1)
17            else:
18                # Deterministic: pick the most likely token
19                _, idx_next = torch.topk(probs, k=1, dim=-1)
20
21            idx = torch.cat((idx, idx_next), dim=1)
22
23    return idx
24
25 # From getting most probable next character
26 def get_best_prediction(logits):
27     probs = F.softmax(logits, dim=-1)
28     best_prob, best_idx = torch.topk(probs, k=1, dim=-1)
29     return best_idx, best_prob

```


Chapter 13

Generative Models

13.1 Generative Adversarial Networks (GANs)

13.1.1 Basic GAN Architecture

Purpose: Generate realistic data through adversarial training between generator and discriminator.

Simple Example - DCGAN:

```
1 import torch.nn as nn
2
3 class Generator(nn.Module):
4     def __init__(self, nz=100, ngf=64, nc=3):
5         super().__init__()
6         # nz: noise dimension, ngf: generator feature maps, nc: channels
7         self.main = nn.Sequential(
8             # Input: (batch, nz, 1, 1)
9             nn.ConvTranspose2d(nz, ngf * 8, 4, 1, 0, bias=False),
10            nn.BatchNorm2d(ngf * 8),
11            nn.ReLU(True),
12
13            # State: (batch, ngf*8, 4, 4)
14            nn.ConvTranspose2d(ngf * 8, ngf * 4, 4, 2, 1, bias=False),
15            nn.BatchNorm2d(ngf * 4),
16            nn.ReLU(True),
17
18            # State: (batch, ngf*4, 8, 8)
19            nn.ConvTranspose2d(ngf * 4, ngf * 2, 4, 2, 1, bias=False),
20            nn.BatchNorm2d(ngf * 2),
21            nn.ReLU(True),
22
23            # State: (batch, ngf*2, 16, 16)
24            nn.ConvTranspose2d(ngf * 2, ngf, 4, 2, 1, bias=False),
25            nn.BatchNorm2d(ngf),
26            nn.ReLU(True),
27
28            # Output: (batch, nc, 32, 32)
```

```

29         nn.ConvTranspose2d(ngf, nc, 4, 2, 1, bias=False),
30         nn.Tanh()
31     )
32
33     def forward(self, x):
34         return self.main(x)
35
36 class Discriminator(nn.Module):
37     def __init__(self, nc=3, ndf=64):
38         super().__init__()
39         # nc: input channels, ndf: discriminator feature maps
40         self.main = nn.Sequential(
41             # Input: (batch, nc, 32, 32)
42             nn.Conv2d(nc, ndf, 4, 2, 1, bias=False),
43             nn.LeakyReLU(0.2, inplace=True),
44
45             # State: (batch, ndf, 16, 16)
46             nn.Conv2d(ndf, ndf * 2, 4, 2, 1, bias=False),
47             nn.BatchNorm2d(ndf * 2),
48             nn.LeakyReLU(0.2, inplace=True),
49
50             # State: (batch, ndf*2, 8, 8)
51             nn.Conv2d(ndf * 2, ndf * 4, 4, 2, 1, bias=False),
52             nn.BatchNorm2d(ndf * 4),
53             nn.LeakyReLU(0.2, inplace=True),
54
55             # State: (batch, ndf*4, 4, 4)
56             nn.Conv2d(ndf * 4, ndf * 8, 4, 2, 1, bias=False),
57             nn.BatchNorm2d(ndf * 8),
58             nn.LeakyReLU(0.2, inplace=True),
59
60             # Output: (batch, 1, 1, 1)
61             nn.Conv2d(ndf * 8, 1, 4, 1, 0, bias=False),
62             nn.Sigmoid()
63         )
64
65     def forward(self, x):
66         return self.main(x).view(-1, 1).squeeze(1)
67
68 # Initialize networks
69 netG = Generator()
70 netD = Discriminator()
71
72 print(f"Generator parameters: {sum(p.numel() for p in netG.parameters())}")
73 print(f"Discriminator parameters: {sum(p.numel() for p in netD.parameters())}")

```

Complex Example - GAN Training Loop:

```

1  def train_gan(netG, netD, dataloader, num_epochs=5, lr=0.0002, beta1=0.5):
2      device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
3      netG.to(device)
4      netD.to(device)
5
6      # Loss function and optimizers
7      criterion = nn.BCELoss()
8      optimizerD = torch.optim.Adam(netD.parameters(), lr=lr, betas=(beta1, 0.999))
9      optimizerG = torch.optim.Adam(netG.parameters(), lr=lr, betas=(beta1, 0.999))
10
11     # Fixed noise for visualization
12     fixed_noise = torch.randn(64, 100, 1, 1, device=device)
13
14     # Labels
15     real_label = 1.0
16     fake_label = 0.0
17
18     for epoch in range(num_epochs):
19         for i, (data, _) in enumerate(dataloader):
20             #####
21             # (1) Update D network: maximize log(D(x)) + log(1 - D(G(z)))
22             #####
23             netD.zero_grad()
24
25             # Train with real batch
26             real_data = data.to(device)
27             batch_size = real_data.size(0)
28             label = torch.full((batch_size,), real_label, dtype=torch.float,
29                               ↪ device=device)
29
30             output = netD(real_data)
31             errD_real = criterion(output, label)
32             errD_real.backward()
33             D_x = output.mean().item()
34
35             # Train with fake batch
36             noise = torch.randn(batch_size, 100, 1, 1, device=device)
37             fake = netG(noise)
38             label.fill_(fake_label)
39             output = netD(fake.detach())
40             errD_fake = criterion(output, label)
41             errD_fake.backward()
42             D_G_z1 = output.mean().item()
43             errD = errD_real + errD_fake
44             optimizerD.step()
45
46             #####

```

```

47         # (2) Update G network: maximize log(D(G(z)))
48         #####
49         netG.zero_grad()
50         label.fill_(real_label) # Fake labels are real for generator cost
51         output = netD(fake)
52         errG = criterion(output, label)
53         errG.backward()
54         D_G_z2 = output.mean().item()
55         optimizerG.step()
56
57         # Print statistics
58         if i % 50 == 0:
59             print(f'[{epoch}/{num_epochs}] [{i}/{len(dataloader)}] '
60                   f'Loss_D: {errD.item():.4f} Loss_G: {errG.item():.4f} '
61                   f'D(x): {D_x:.4f} D(G(z)): {D_G_z1:.4f} / {D_G_z2:.4f}')
62
63         # Generate images for visualization
64         with torch.no_grad():
65             fake_images = netG(fixed_noise)
66             # Save or display fake_images here
67
68     return netG, netD
69
70 # Weight initialization (important for GAN training)
71 def weights_init(m):
72     classname = m.__class__.__name__
73     if classname.find('Conv') != -1:
74         nn.init.normal_(m.weight.data, 0.0, 0.02)
75     elif classname.find('BatchNorm') != -1:
76         nn.init.normal_(m.weight.data, 1.0, 0.02)
77         nn.init.constant_(m.bias.data, 0)
78
79 # Apply weight initialization
80 netG.apply(weights_init)
81 netD.apply(weights_init)

```

13.2 Variational Autoencoders (VAEs)

13.2.1 VAE Architecture

Purpose: Learn probabilistic latent representations for generation and reconstruction.

Complex Example - VAE Implementation:

```

1 class VAE(nn.Module):
2     def __init__(self, input_dim=784, hidden_dim=400, latent_dim=20):
3         super().__init__()
4

```

```

5         # Encoder
6         self.encoder = nn.Sequential(
7             nn.Linear(input_dim, hidden_dim),
8             nn.ReLU(),
9             nn.Linear(hidden_dim, hidden_dim),
10            nn.ReLU()
11        )
12
13        # Latent space parameters
14        self.fc_mu = nn.Linear(hidden_dim, latent_dim)
15        self.fc_logvar = nn.Linear(hidden_dim, latent_dim)
16
17        # Decoder
18        self.decoder = nn.Sequential(
19            nn.Linear(latent_dim, hidden_dim),
20            nn.ReLU(),
21            nn.Linear(hidden_dim, hidden_dim),
22            nn.ReLU(),
23            nn.Linear(hidden_dim, input_dim),
24            nn.Sigmoid() # For image data normalized to [0,1]
25        )
26
27        def encode(self, x):
28            h = self.encoder(x)
29            mu = self.fc_mu(h)
30            logvar = self.fc_logvar(h)
31            return mu, logvar
32
33        def reparameterize(self, mu, logvar):
34            std = torch.exp(0.5 * logvar)
35            eps = torch.randn_like(std)
36            return mu + eps * std
37
38        def decode(self, z):
39            return self.decoder(z)
40
41        def forward(self, x):
42            mu, logvar = self.encode(x)
43            z = self.reparameterize(mu, logvar)
44            recon_x = self.decode(z)
45            return recon_x, mu, logvar
46
47        # VAE Loss function
48        def vae_loss(recon_x, x, mu, logvar, beta=1.0):
49            # Reconstruction loss (binary cross entropy)
50            BCE = F.binary_cross_entropy(recon_x, x, reduction='sum')
51
52            # KL divergence loss

```

```

53     KLD = -0.5 * torch.sum(1 + logvar - mu.pow(2) - logvar.exp())
54
55     # Beta-VAE: beta controls the weight of KL divergence
56     return BCE + beta * KLD
57
58 # Training function
59 def train_vae(model, dataloader, epochs=10, lr=1e-3, beta=1.0):
60     optimizer = torch.optim.Adam(model.parameters(), lr=lr)
61     model.train()
62
63     for epoch in range(epochs):
64         train_loss = 0
65         for batch_idx, (data, _) in enumerate(dataloader):
66             data = data.view(-1, 784) # Flatten images
67             optimizer.zero_grad()
68
69             # Forward pass
70             recon_batch, mu, logvar = model(data)
71
72             # Compute loss
73             loss = vae_loss(recon_batch, data, mu, logvar, beta)
74
75             # Backward pass
76             loss.backward()
77             train_loss += loss.item()
78             optimizer.step()
79
80             if batch_idx % 100 == 0:
81                 print(f'Epoch: {epoch}, Batch: {batch_idx}, '
82                       f'Loss: {loss.item() / len(data):.6f}')
83
84         print(f'Epoch: {epoch}, Average loss: {train_loss /
85               ↪ len(dataloader.dataset):.6f}')
86
87 # Generate new samples
88 @torch.no_grad()
89 def generate_samples(model, num_samples=64, latent_dim=20):
90     model.eval()
91     z = torch.randn(num_samples, latent_dim)
92     samples = model.decode(z)
93     return samples
94
95 # Usage
96 vae = VAE(input_dim=784, hidden_dim=400, latent_dim=20)
97 # Assuming you have a dataloader for MNIST or similar
98 # train_vae(vae, train_dataloader)
99 # generated_images = generate_samples(vae)

```


13.3 Advanced GAN Variants

13.3.1 Conditional GAN (cGAN)

Purpose: Generate data conditioned on class labels or other information.

Example - Conditional Generator:

```

1  class ConditionalGenerator(nn.Module):
2      def __init__(self, num_classes=10, nz=100, ngf=64, nc=3):
3          super().__init__()
4          self.num_classes = num_classes
5
6          # Embedding for class labels
7          self.label_embed = nn.Embedding(num_classes, num_classes)
8
9          # Main generator network
10         self.main = nn.Sequential(
11             # Input: noise + label embedding
12             nn.ConvTranspose2d(nz + num_classes, ngf * 8, 4, 1, 0, bias=False),
13             nn.BatchNorm2d(ngf * 8),
14             nn.ReLU(True),
15
16             nn.ConvTranspose2d(ngf * 8, ngf * 4, 4, 2, 1, bias=False),
17             nn.BatchNorm2d(ngf * 4),
18             nn.ReLU(True),
19
20             nn.ConvTranspose2d(ngf * 4, ngf * 2, 4, 2, 1, bias=False),
21             nn.BatchNorm2d(ngf * 2),
22             nn.ReLU(True),
23
24             nn.ConvTranspose2d(ngf * 2, ngf, 4, 2, 1, bias=False),
25             nn.BatchNorm2d(ngf),
26             nn.ReLU(True),
27
28             nn.ConvTranspose2d(ngf, nc, 4, 2, 1, bias=False),
29             nn.Tanh()
30         )
31
32     def forward(self, noise, labels):
33         # Embed labels and reshape
34         label_embed = self.label_embed(labels).view(labels.size(0),
35             ↪ self.num_classes, 1, 1)
36
37         # Concatenate noise and label embedding
38         gen_input = torch.cat([noise, label_embed], dim=1)
39
40         return self.main(gen_input)
41
42     # Generate specific classes

```

```
42 def generate_class_samples(model, class_label, num_samples=16):
43     model.eval()
44     with torch.no_grad():
45         noise = torch.randn(num_samples, 100, 1, 1)
46         labels = torch.full((num_samples,), class_label, dtype=torch.long)
47         generated = model(noise, labels)
48     return generated
49
50 # Usage
51 cgan_generator = ConditionalGenerator(num_classes=10)
52 # Generate samples of class 7
53 class_7_samples = generate_class_samples(cgan_generator, class_label=7)
```

Chapter 14

Complete Examples and Applications

14.1 Building a Simple Neural Network

Complete MLP Example:

```
1  import torch
2  import torch.nn as nn
3  import torch.nn.functional as F
4  import torch.optim as optim
5
6  class SimpleMLP(nn.Module):
7      def __init__(self, input_size, hidden_size, output_size):
8          super(SimpleMLP, self).__init__()
9          self.fc1 = nn.Linear(input_size, hidden_size)
10         self.fc2 = nn.Linear(hidden_size, hidden_size)
11         self.fc3 = nn.Linear(hidden_size, output_size)
12
13         def forward(self, x):
14             x = F.relu(self.fc1(x))
15             x = F.relu(self.fc2(x))
16             x = self.fc3(x)
17             return x
18
19         # Training function
20         def train_mlp():
21             # Create model, data, optimizer
22             model = SimpleMLP(784, 128, 10) # MNIST-like dimensions
23             optimizer = optim.AdamW(model.parameters(), lr=0.001)
24
25             # Training loop
26             for epoch in range(100):
27                 # Generate dummy batch
28                 x = torch.randn(32, 784)
29                 y = torch.randint(0, 10, (32,))
30
```

```
31     # Forward pass
32     logits = model(x)
33     loss = F.cross_entropy(logits, y)
34
35     # Backward pass
36     optimizer.zero_grad()
37     loss.backward()
38     optimizer.step()
39
40     if epoch % 20 == 0:
41         print(f'Epoch {epoch}, Loss: {loss.item():.4f}')
42
43 if __name__ == "__main__":
44     train_mlp()
```

14.2 Character-Level Language Model

Complete Bigram Model Example:

```
1  import torch
2  import torch.nn.functional as F
3
4  class BigramLanguageModel:
5      def __init__(self, vocab_size):
6          self.vocab_size = vocab_size
7          # Initialize weight matrix
8          g = torch.Generator().manual_seed(2147483647)
9          self.W = torch.randn((vocab_size, vocab_size),
10                               generator=g, requires_grad=True)
11
12     def forward(self, xs, ys):
13         # Convert to one-hot
14         xenc = F.one_hot(xs, num_classes=self.vocab_size).float()
15
16         # Neural network forward pass
17         logits = xenc @ self.W
18         counts = logits.exp()
19         probs = counts / counts.sum(1, keepdims=True)
20
21         # Compute loss
22         loss = -probs[torch.arange(len(ys)), ys].log().mean()
23         return loss, probs
24
25     def generate(self, num_samples=5):
26         g = torch.Generator().manual_seed(2147483647)
27
```

```

28     for i in range(num_samples):
29         out = []
30         ix = 0 # Start token
31
32         while True:
33             # Get probabilities for current character
34             xenc = F.one_hot(torch.tensor([ix]),
35                               num_classes=self.vocab_size).float()
36             logits = xenc @ self.W
37             counts = logits.exp()
38             probs = counts / counts.sum(1, keepdims=True)
39
40             # Sample next character
41             ix = torch.multinomial(probs, num_samples=1,
42                                   replacement=True, generator=g).item()
43
44             out.append(ix)
45
46             if ix == 0: # Stop token
47                 break
48
49         yield out
50
51 def train(self, xs, ys, learning_rate=50, num_steps=100):
52     for step in range(num_steps):
53         # Forward pass
54         loss, probs = self.forward(xs, ys)
55
56         # Backward pass
57         self.W.grad = None
58         loss.backward()
59
60         # Update weights
61         self.W.data += -learning_rate * self.W.grad
62
63         if step % 20 == 0:
64             print(f'Step {step}, Loss: {loss.item():.4f}')
65
66 # Usage example
67 if __name__ == "__main__":
68     # Create dummy data (character indices)
69     vocab_size = 27
70     xs = torch.tensor([0, 5, 13, 13, 1]) # Input characters
71     ys = torch.tensor([5, 13, 13, 1, 0]) # Target characters
72
73     # Create and train model
74     model = BigramLanguageModel(vocab_size)
75     model.train(xs, ys)

```

```

76     # Generate samples
77     print("Generated sequences:")
78     for i, sequence in enumerate(model.generate(3)):
79         print(f"Sample {i+1}: {sequence}")

```

14.3 Transformer Language Model Excerpt

Key Components from Educational Materials:

```

1  class CausalSelfAttention(nn.Module):
2      def __init__(self, config):
3          super().__init__()
4          assert config.n_embd % config.n_head == 0
5
6          # Key, query, value projections for all heads
7          self.c_attn = nn.Linear(config.n_embd, 3 * config.n_embd)
8          self.c_proj = nn.Linear(config.n_embd, config.n_embd)
9
10         # Causal mask
11         self.register_buffer("bias",
12                               torch.tril(torch.ones(config.block_size, config.block_size))
13                               .view(1, 1, config.block_size, config.block_size))
14
15         self.n_head = config.n_head
16         self.n_embd = config.n_embd
17
18     def forward(self, x):
19         B, T, C = x.size()
20
21         # Calculate query, key, values for all heads
22         q, k, v = self.c_attn(x).split(self.n_embd, dim=2)
23
24         # Reshape for multi-head attention
25         k = k.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
26         q = q.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
27         v = v.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
28
29         # Causal self-attention
30         att = (q @ k.transpose(-2, -1)) * (1.0 / math.sqrt(k.size(-1)))
31         att = att.masked_fill(self.bias[:, :, :T, :T] == 0, float('-inf'))
32         att = F.softmax(att, dim=-1)
33         y = att @ v
34
35         # Re-assemble all head outputs
36         y = y.transpose(1, 2).contiguous().view(B, T, C)
37         y = self.c_proj(y)

```

```
38
39     return y
40
41 class Block(nn.Module):
42     def __init__(self, config):
43         super().__init__()
44         self.ln_1 = nn.LayerNorm(config.n_embd)
45         self.attn = CausalSelfAttention(config)
46         self.ln_2 = nn.LayerNorm(config.n_embd)
47         self.mlp = nn.ModuleDict(dict(
48             c_fc    = nn.Linear(config.n_embd, 4 * config.n_embd),
49             c_proj   = nn.Linear(4 * config.n_embd, config.n_embd),
50             act      = nn.GELU(),
51         ))
52
53     def forward(self, x):
54         x = x + self.attn(self.ln_1(x))
55         x = x + self.mlp.c_proj(self.mlp.act(self.mlp.c_fc(self.ln_2(x))))
56         return x
```


Chapter 15

PyTorch 2.x Features and Optimizations

15.1 torch.compile for Performance

Purpose: Compile PyTorch models for significant performance improvements.

Simple Example:

```
1  import torch
2  import torch.nn as nn
3
4  # Simple model
5  class SimpleModel(nn.Module):
6      def __init__(self):
7          super().__init__()
8          self.linear = nn.Linear(10, 1)
9
10     def forward(self, x):
11         return self.linear(x)
12
13     # Regular model
14     model = SimpleModel()
15
16     # Compiled model - faster execution
17     compiled_model = torch.compile(model)
18
19     # Use compiled model
20     x = torch.randn(32, 10)
21     output = compiled_model(x) # Faster than model(x)
```

Complex Example:

```
1  # Compile with different backends and modes
2  model = torch.compile(model, backend="inductor", mode="max-autotune")
3
4  # For inference only
```

```
5 model = torch.compile(model, mode="reduce-overhead")
6
7 # Full graph compilation
8 model = torch.compile(model, fullgraph=True)
9
10 # Training loop with compiled model
11 compiled_model = torch.compile(model)
12 for batch in dataloader:
13     optimizer.zero_grad()
14     loss = compiled_model(batch.x, batch.y)
15     loss.backward()
16     optimizer.step()
```

15.2 Mixed Precision Training

Purpose: Use automatic mixed precision for faster training with lower memory usage.

Simple Example:

```
1 import torch
2 from torch.cuda.amp import autocast, GradScaler
3
4 # Create scaler for gradient scaling
5 scaler = GradScaler()
6
7 # Training loop with mixed precision
8 for batch in dataloader:
9     optimizer.zero_grad()
10
11     # Forward pass with autocast
12     with autocast():
13         outputs = model(inputs)
14         loss = criterion(outputs, targets)
15
16     # Scaled backward pass
17     scaler.scale(loss).backward()
18     scaler.step(optimizer)
19     scaler.update()
```

Complex Example:

```
1 # Device-agnostic autocast (PyTorch 2.x)
2 device_type = "cuda" if torch.cuda.is_available() else "cpu"
3
4 for batch in dataloader:
5     with torch.autocast(device_type=device_type, dtype=torch.float16):
6         logits = model(batch.input_ids)
```

```
7         loss = F.cross_entropy(logits, batch.labels)
8
9         if device_type == "cuda":
10             scaler.scale(loss).backward()
11             scaler.unscale_(optimizer)
12             torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
13             scaler.step(optimizer)
14             scaler.update()
15         else:
16             loss.backward()
17             optimizer.step()
```

15.3 Improved DataLoader and Data Handling

Purpose: Use new PyTorch 2.x data loading features for better performance.

Simple Example:

```
1 from torch.utils.data import DataLoader
2
3 # PyTorch 2.x: Better multiprocessing
4 dataloader = DataLoader(
5     dataset,
6     batch_size=32,
7     num_workers=4,
8     persistent_workers=True, # Keep workers alive between epochs
9     pin_memory=True,
10    prefetch_factor=2 # Prefetch batches
11 )
12
13 # Non-blocking transfer
14 for batch in dataloader:
15     inputs = batch[0].to(device, non_blocking=True)
16     targets = batch[1].to(device, non_blocking=True)
```

15.4 Better Device and Dtype Handling

Purpose: Use improved PyTorch 2.x device and dtype management.

Example:

```
1 # Direct device and dtype specification
2 device = "cuda" if torch.cuda.is_available() else "cpu"
3
4 # Create tensors directly on device
5 x = torch.randn(100, 10, device=device, dtype=torch.float16)
6
```

```
7  # Model initialization with device/dtype
8  model = nn.Linear(10, 1, device=device, dtype=torch.float32)
9
10 # Tensor creation with factory functions
11 zeros_gpu = torch.zeros(10, 10, device=device)
12 ones_gpu = torch.ones_like(zeros_gpu)
13
14 # Better context managers
15 with torch.device(device):
16     temp_tensor = torch.randn(5, 5) # Automatically on device
```

Chapter 16

Best Practices and Common Patterns

16.1 Memory Management

Efficient Gradient Handling:

```
1  # Clear gradients efficiently
2  model.zero_grad(set_to_none=True) # More memory efficient than zero_grad()
3
4  # Use torch.no_grad() for inference
5  @torch.no_grad()
6  def evaluate_model(model, data_loader):
7      model.eval()
8      total_loss = 0
9      for batch in data_loader:
10         outputs = model(batch)
11         loss = compute_loss(outputs, batch.targets)
12         total_loss += loss.item()
13     return total_loss / len(data_loader)
14
15 # PyTorch 2.x: Use torch.inference_mode() for even better performance
16 @torch.inference_mode()
17 def fast_inference(model, x):
18     return model(x)
19
20 # PyTorch 2.x: Compiled inference for maximum speed
21 @torch.compile
22 @torch.inference_mode()
23 def compiled_inference(model, x):
24     return model(x)
```

16.2 Device Management

GPU/CPU Handling:

```

1  # PyTorch 2.x: Better device handling
2  device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
3  print(f"Using device: {device}")
4
5  # PyTorch 2.x: Direct device specification in tensor creation
6  data = torch.randn(10, 3, device=device)
7  model = model.to(device)
8
9  # PyTorch 2.x: Context manager for device
10 with torch.device(device):
11     x = torch.randn(5, 3) # Automatically on the specified device
12
13 # In training loop
14 for batch in data_loader:
15     # Move batch to device
16     batch = [t.to(device, non_blocking=True) for t in batch] # non_blocking for
    ↪ speed
17     X, Y = batch
18
19     # Forward pass
20     logits, loss = model(X, Y)
21
22     # Synchronize for accurate timing (CUDA only)
23     if device.type == 'cuda':
24         torch.cuda.synchronize()

```

16.3 Common Debugging Techniques

Shape and Gradient Debugging:

```

1  # Monitor tensor shapes
2  def debug_shapes(x, name="tensor"):
3      print(f"{name} shape: {x.shape}, dtype: {x.dtype}, device: {x.device}")
4      if x.requires_grad:
5          print(f"{name} requires grad: {x.requires_grad}")
6      return x
7
8  # Check gradients
9  def check_gradients(model):
10     for name, param in model.named_parameters():
11         if param.grad is not None:
12             grad_norm = param.grad.norm()
13             print(f"{name}: grad_norm = {grad_norm:.6f}")
14         else:
15             print(f"{name}: no gradient computed")
16

```

```
17 # Monitor loss and learning
18 def training_step_with_monitoring(model, optimizer, batch):
19     X, Y = batch
20
21     # Forward pass
22     logits, loss = model(X, Y)
23
24     # Check for NaN
25     if torch.isnan(loss):
26         print("WARNING: NaN loss detected!")
27         return
28
29     # Backward pass
30     model.zero_grad(set_to_none=True)
31     loss.backward()
32
33     # Check gradient norms
34     total_grad_norm = 0
35     for param in model.parameters():
36         if param.grad is not None:
37             total_grad_norm += param.grad.norm().item() ** 2
38     total_grad_norm = total_grad_norm ** 0.5
39
40     print(f"Loss: {loss.item():.6f}, Grad norm: {total_grad_norm:.6f}")
41
42     optimizer.step()
```


Chapter 17

Conclusion

This tutorial has covered PyTorch functions from fundamental tensor operations to complete neural network implementations. The progression from basic operations like `torch.tensor()` and `torch.zeros()` to complex architectures like Transformers demonstrates how these building blocks combine to create powerful machine learning models.

Key takeaways:

- Start with tensor fundamentals before moving to neural networks
- Understand shapes and broadcasting for effective debugging
- Use automatic differentiation properly with `requires_grad`
- Master the core operations: matrix multiplication, softmax, cross-entropy
- Practice with complete examples to solidify understanding
- Follow best practices for memory and device management

The examples drawn from educational materials and neural network implementations provide real-world context for how these functions are used in practice. Continue practicing with these patterns and gradually build more complex models.

Further Reading

- PyTorch Official Documentation: <https://pytorch.org/docs/>
- PyTorch Tutorials: <https://pytorch.org/tutorials/>
- Deep Learning with PyTorch: <https://pytorch.org/deep-learning-with-pytorch>